

All Paths Lead to the Root

Théophile Brézot ^{*} Chloé Hébant [†]

Abstract

In an attempt to fix the defects of the definition of forward security for Symmetric Searchable Encryption (SSE) schemes, Amjad et al. [2] proposed injection security. This new security property is strictly stronger than most security properties known to date, which makes it particularly challenging to design schemes meeting its requirements. In this work, we show how it is possible to use trees to decorrelate the modification of an index from its effects, hence achieving injection security. In addition to being conceptually simple, our scheme features non-interactive, stateless and mutation-free search operations that allow supporting concurrent readers easily. Finally, the proposed reference implementation is efficient: both Insert and Search operations execute in milliseconds even when operating on an index with up to a million entries and volumes up to a thousand.

1 Introduction

Symmetric Searchable Encryption (SSE) is a formalism used to study the security of protocols manipulating a document index stored by a distant server: an SSE scheme is a bipartite protocol with one of the party (the server) usually being considered honest-but-curious, and the notion of leakage is used to capture the amount of information that party can acquire by observing the execution of SSE operations. There exists a natural tension between the efficiency of a scheme and its security. SSE schemes resolutely fall into the practical side when compared to more secure schemes like ORAM, but successive innovations have led to constructions proposing stronger security guarantees to the detriment of performance. We stress that in many real-world applications, the performance of a scheme is almost as important as its security since poor performances may hinder its adoption, resulting in essentially no security. While server-side memory is generally not an issue, client-side memory is a particularly constrained resource. However, apart from notable works like [11], most SSE schemes have a space complexity that is linear in the number of indexed labels L . Another source of inefficiency that runs even deeper into the SSE literature is the impossibility to securely perform concurrent queries to the index, which stems from the presence of client-side mutable state and the lack of a clear security model. Some recent developments in [1, 7] made a step in the direction of formalizing security in the presence of concurrent queries, but to the best of our knowledge, no concurrent SSE scheme with a decent security has ever been published. In this work, we propose an SSE scheme with non-interactive, stateless and

^{*}Cosmian: theophile.brezot@cosmian.com

[†]Cosmian: chloe.hebant@cosmian.com

mutation-free search operations that make it trivially amenable to concurrent search queries provided the server guarantees basic shared-object properties. We believe our scheme opens up a broad range of applications in which lightweight clients query an index managed by one powerful client.

On the security front, forward security – formalized by Bost in [4] as the unlinkability of updates with previous operations – has been considered as a state-of-the-art security property, even though it does not guarantee that future search queries cannot be linked to past updates. A persistent adversary thus eventually learns what forward security attempts to hide. Injection Security is a much stronger property – defined as the unlinkability of updates with all other past and future operations – that has recently been proposed [2] to remedy forward security’s shortcomings. We believe it to be a real step forward, which is why we present in this work a new injection-secure SSE scheme named PLOC.

1.1 Related works

The first and only injection-secure SSE scheme to date [2] relies on decorrelating the server-side effects of a modification from its intent as follows: upon modifying bindings for a label ℓ , the client saves this modification locally under this label and performs instead the modifications saved under a different label ℓ' chosen independently from ℓ . In their scheme the authors used a round-robin strategy going through all the labels by batch. This implies that labels are updated server-side every $\frac{|\mathbb{L}|}{C}$ modifications, where \mathbb{L} is the labeling domain and C the batch size. Conversely, the size of the client-side state is $O(\frac{|\mathbb{L}|}{C})$ in the worst case. Moreover, in order to achieve injection security, the client must modify a number of values that is independent from ℓ , which results in a bandwidth in $O(C \times S)$ for some constant S chosen at setup. Conscious of this cost, the authors astutely pad and truncate the set of values indexed under each label to enable users to choose $S = (1 - \epsilon) \cdot V$ for some loss parameter $0 < \epsilon \leq 1$ and V an upper-bound on the number of values indexed per label, also chosen at setup. However, the loss cannot be increased past a certain amount without severely damaging the search results, and the bandwidth cost remains $O(C \times V)$. The users are therefore caught in an impossible dilemma: they must constrain the size of the labeling domain to limit both the client-side stash and insertion bandwidth at the risk of speeding-up query-recovery attacks (since possible labels are fewer, it is easier to tell them apart), or otherwise choose between an acceptable bound on the client-side storage or the bandwidth. Finally, since the very security of this scheme relies on delaying the modification of the server state, it seems impossible to use it as a basis on which to develop a concurrent SSE scheme with anything but a very relaxed concurrency model. All in all, this scheme represents a break-through in the field, but presents costs that may be prohibitively high for many use-cases.

Even though not achieving injection security (nor forward security), we quickly present here *2ch* [3] since it is conceptually close to ours. This scheme is based on a complete binary tree, in which indexing a value under a label consists in deriving two candidate branches from this label and its current volume using a PRF. The client then fetches those two branches and chooses to place the new value inside the least filled one. Search operations always fetch $2V$ branches which first $2\nu(\ell)$ correspond to the branches derived to bind values

to this label so far, where $\nu(\ell)$ is the current volume of the label ℓ . Using this technique allows implementing a volume-hiding [15] scheme. The authors also managed to reduce the bandwidth from $O(\log B)$ to $O(\log \log B)$ where B is the maximum number of bindings that can be indexed in the tree by beheading the tree to create a forest of trees of depth $\log \log B$. The subsequent schemes proposed in the same article build on and attempt to recover the forward and backward securities [6], but are sub-optimal since they require search operations to perform more work.

1.2 Contributions

In this work, we propose a novel injection-secure SSE scheme based on a different technique to decorrelate a mutation intent from its effects. Like in [3], we rely on a tree structure and values are stored in some node alongside the path leading to the leaf selected using a PRF of the indexing label and its current volume. However, since all paths lead to the root of a tree, the client can fetch truly random branches upon insertion: as long as it possible to compact the resulting subtree enough to make space at the root, the new value can be inserted on its target path. Upon search, all V branches associated to a label are read in order to find the indexed values. Using this technique allows us to build a simple SSE scheme with very desirable properties:

- insertions are only required to maintain a $O(L)$ client-side stash, where L is the number of labels currently indexed, to store their associated volumes. This is a significant improvement over the $O(\frac{|L|}{C})$ bound from [2]. Our scheme implements insert operations using a single round-trip and $O(cn (\lg(\frac{B}{n}) + 2))$ bandwidth, where n is the number of branches fetched from the tree for compaction, c the capacity of each node and B the total number of indexed bindings;
- searches need only know the key and do not perform any mutation, which allows them to be trivially concurrent and hence to implement a single-writer, several-readers system that only leaks the search pattern. Even though our insertions cannot be used by concurrent clients due to the mutable state that is the label-volume map, we consider it a net improvement over the scheme proposed in the recent article from Agarwal et al. [1] that doesn't even guarantee forward security in the context of concurrent operations. Our scheme implements search operations using a single round-trip and a $O(c \cdot V \lg B)$ bandwidth.
- additionally, we show how to apply some of the standard techniques from the literature to implement a more expressive interface, for example that support deletions, and to make insertions stateless.

2 Preliminaries

2.1 Structured Encryption

Structured Encryption (STE) was first introduced in [8] to allow a client to query a data structure stored by a server with controlled leakages. An STE scheme is based on an Abstract Data Type (ADT) defined as follows:

Definition 2.1 (Abstract Data Type). An abstract data type is defined as a pair $(\mathbb{S}, \mathcal{O})$ of states and operations on those states which every realizations of this abstract data type must support. Without loss of generality, we constrain operations in \mathcal{O} to be parameterized by their argument and result sets \mathbb{A} and \mathbb{R} and to be of type $\mathbb{A} \times \mathbb{S} \rightarrow \mathbb{R} \times \mathbb{S}$.

Definition 2.2 (Structured Encryption Scheme). Given an abstract data type $\mathcal{T} = (\mathbb{S}, \mathcal{O})$, a \mathcal{T} encryption scheme is a bipartite protocol involving a client and a server that supports the following operations:

- $(k, \text{st}; EDS) \leftarrow \text{Setup}(DS;)$ in which the client takes as input a valid data structure $DS \in \mathbb{S}$, and at the end of which the client receives a key k and a state st , and the server receives the encrypted data structure EDS corresponding to DS ;

for each operation $\text{op} \in \mathcal{O}$ parameterized by \mathbb{A} and \mathbb{R} :

- $(\text{res}, \text{st}'; EDS') \leftarrow \text{op}(k, \text{args}, \text{st}; EDS)$ in which the client takes as input the key k , a state st and some arguments $\text{args} \in \mathbb{A}$, the server takes as input an encrypted data structure EDS corresponding to a data-structure $DS \in \mathbb{S}$, and at the end of which the client receives the results $\text{res} \in \mathbb{R}$ and the server the encrypted data structure EDS' corresponding to the data structure $DS' \in \mathbb{S}$, where res and DS' correspond to what locally applying op on args and DS would return.

Remark. Even though the standard definition for a structured encryption scheme expects the setup to take as argument an initial data structure, the schemes described in this article do not. However, it is easy to see that an alternative setup procedure taking as argument a data structure can be obtained without additional leakage by executing a setup followed by one or more update operations. Indeed, PLOC insert operation only has trivial leakages and so does the tree merge operation when taking the entire tree as argument.

2.2 Multi-map Encryption Scheme

One particular instance of structured encryption has been at the center of most of the works in the literature for its immediate applicability: the multi-map encryption scheme.

Definition 2.3 (Multi-Map). A (semi-dynamic) multi-map is an abstract data type parameterized by a label domain \mathbb{L} and a value domain \mathbb{V} , which valid states are the functions in $\mathbb{L} \rightarrow \mathbb{V}^*$, where $\mathbb{V}^* = \perp \cup \mathbb{V} \times \mathbb{V}^*$, and which operations are:

- $((v_i)_i, MM) \leftarrow \text{Search}(\ell, MM)$ takes as argument a label ℓ and a valid state MM , and returns the (possibly empty) list of values bound to that label by the given state;
- $(\perp, MM') \leftarrow \text{Insert}((\ell, v), MM)$ takes as argument a label $\ell \in \mathbb{L}$, a value $v \in \mathbb{V}$ and a state MM , and returns the new state MM' such that $MM'(\ell') = MM(\ell')$ for each label $\ell' \neq \ell$ and $MM'(\ell) = (v, MM(\ell))$ otherwise.

A multi-map is said to be fully-dynamic if it additionally supports:

- $(\perp, \text{MM}') \leftarrow \text{Delete}((\ell, v), \text{MM})$ takes as argument a label $\ell \in \mathbb{L}$, a value $v \in \mathbb{V}$ and a state MM , and returns the new state MM' such that $\text{MM}'(\ell') = \text{MM}(\ell')$ for each label $\ell' \neq \ell$ and $\text{MM}'(\ell) = \text{MM}(\ell) \setminus \ell$ otherwise.

Definition 2.4 (Multi-Map Encryption Scheme). A (semi-dynamic) multi-map encryption scheme is therefore the protocol exposing the following operations which semantics is given by the STE definition:

- $((v_i)_i, \text{st}'; \text{EMM}) \leftarrow \text{Search}(k, \ell, \text{st}; \text{EMM});$
- $(\perp, \text{st}'; \text{EMM}') \leftarrow \text{Insert}(k, (\ell, v), \text{st}; \text{EMM}).$

A multi-map encryption scheme is said to be fully-dynamic if it additionally exposes:

- $(\perp, \text{st}'; \text{EMM}') \leftarrow \text{Delete}(k, (\ell, v), \text{st}; \text{EMM}).$

2.3 Symmetric Searchable Encryption

While structured encryption addresses the abstract and fundamental problem of securing data-structures, Symmetric Searchable Encryption (SSE) as described in [10] by Curtmola et al. addresses the problem of securing a specific use-case in which a client wishes to retrieve the list of document IDs of all documents matching a given label, from a collection of documents stored in a remote server. While the interface defined in [10] differs from those of a multi-map, it is trivial to implement an SSE scheme from a multi-map encryption scheme. For this reason, we claim in this work to present a new SSE scheme while only presenting a new encrypted multi-map implementation.

2.4 Leakages and adaptive security

The leakage function plays a central role in the SSE literature since it is used to quantify the information leaked to the adversary by each operation. In this work, we consider the notion of adaptive security against a honest-but-curious, polynomial-time server.

Definition 2.5 (Adaptive Security). An STE scheme Σ is said to be \mathcal{L} -adaptively-secure if there exists a simulator ϕ such that for any stateful honest-but-curious PPT adversary \mathcal{A} and any auxiliary information $z \in \{0, 1\}^*$:

$$\left| P(\mathbf{Real}_{\Sigma, \mathcal{A}}(1^\lambda, z) = 1) - P(\mathbf{Ideal}_{\Sigma, \mathcal{A}}^{\phi, \mathcal{L}}(1^\lambda, z) = 1) \right| = \mathbf{negl}(\lambda),$$

where the real and ideal games are defined as follows:

- the challenger either executes the client code of the STE scheme Σ (real game) or the composition $\phi \circ \mathcal{L}$ of the simulator with the leakage function (ideal game), while the adversary plays the server code of the STE scheme Σ .
- the adversary adaptively chooses a sequence of (polynomially many in λ) ADT operations that are jointly executed with the challenger.
- the adversary outputs a bit $b \in \{0, 1\}$.

2.5 Security properties

If leakage functions capture the information that is leaked by SE operations, evaluating the level of privacy guaranteed by a given leakage or even establishing an order relation on the leakages that reflects those security guarantees remains a difficult problem [5, 14]. It has thus become idiomatic to describe scheme privacy using *security properties*, which are associated to the set of attacks known to hold against them.

The most important source of attacks against SSE schemes is the search pattern – sometimes called query equality – which allows an attacker to partition a sequence of search queries with respect to the searched label (without always being able to guess which are those labels).

Definition 2.6 (Search Pattern). Given a sequence of m search operations (o_1, \dots, o_m) , the search pattern, noted sp , is the matrix $M \in \{0, 1\}^{m \times m}$ such that $M[i][j] = i$ if and only if the operations o_i and o_j both search for the same label.

Remark. If the label domain \mathbb{L} is countable, there exists a bijection idx from $[0, |\mathbb{L}|[$ to \mathbb{L} and the matrix-based definition of the search pattern is equivalent to leaking $\pi_{[0, |\mathbb{L}|[} \circ idx^{-1}(\ell)$, where $\pi_{[0, |\mathbb{L}|[}$ is a random permutation on $[0, |\mathbb{L}|[$. This definition is easier to use in practice since it is comes up more naturally in the proofs.

Injection security is a property of the leakage functions of multi-map encryption schemes that has recently been defined in [2] guaranteeing that no link can ever be made between a mutation (insertion or deletion) and any other past or future operation. While we refer the reader to the original article for its game-based definition, we give here the sufficient condition we use in this article:

Theorem 1 (Injection Security). *A leakage function \mathcal{L} is said to be injection-secure if $\mathcal{L}(op, args) = (op, sp)$ if op is Search and op otherwise.*

Proof. See the proof of Theorem 3.3 from [2]. □

2.6 Notations

In the rest of this article, we make use of the notations listed in table 1.

3 PLOC

The complete algorithm is given in appendix D.5¹. We have chosen to implement it in the Scheme programming language [12] for its well-defined semantics and conciseness², which incidentally provides a reference implementation for PLOC. A description of the syntax required to understand this code is provided in appendix B. The present section details the most important aspects of the algorithm step by step. It is assumed that the following names are in scope:

- c the capacity of each node (i.e. the maximum number of datums they can store);

¹It is also available as a Github repository.

²The entire reference implementation given in appendix D is less than 400 lines long.

Table 1: Notations used in this article.

notation	description
s	node size
c	node capacity
ℓ	label
v	value
d	depth
n	number of branches selected for compaction
B	upper-bound on the number of bindings
L	upper-bound on the number of labels
N	upper-bound on the number of tree leaves
V	upper-bound on the maximum label volume
ϕ	(s, D) -tree-encryption-scheme simulator
ψ	PLOC simulator

- n the (64-bit) number of branches to compact per insertion;
- B the (64-bit) maximum number of bindings allowed;
- V the (64-bit) maximum volume allowed;
- (rng) a cryptographic random-number generator producing 64-bit values;
- $(H \text{ key label volume})$ a cryptographic hash function taking as input a 128-bit key, a label bit-string of arbitrary size and an integer, and produces a 128-bit integer. Its only required security guarantee is to be indistinguishable from a pseudo-random function;

In order to store and encrypt its data, PLOC relies on an (s, D) -binary-tree encryption scheme which, for the sake of this article, is based on the abstract data type defined as follows:

Definition 3.1 $((s, D)$ Binary Tree). An (s, D) binary tree is an abstract data type parameterized by a node size s and a depth D , whose valid states are the binary trees of depth D with nodes that are isomorphic to s -long byte strings (i.e. tree nodes are serializable to s -long byte strings), and that exposes the following operations:

- $\text{Select}(\{b_i\})$ that returns the subtree defined by the branches $\{b_i\}$;
- $\text{Mutate}(\{b_i\}, tx)$ that mutates the subtree defined by the branches $\{b_i\}$ by replacing it by its image by the transformation tx on it, which should be a function that takes as argument a tree and returns a tree of the same shape.

Where a *branch* is defined as a bit string representing the choices to be made at each level between the left and right children of a node. Such bit strings are encoded as integers in which the choice made at the root is encoded by the bit of exponent zero. This choice of representation allows for efficient and simple tree traversals. The number of bits in a branch needs not be smaller than the tree depth since the additional bits can simply be ignored by the traversal. This allows using the output of H as branch.

An (s, D) -binary-tree encryption scheme therefore exposes the following operations:

- `(tree-setup)` creates a server-side encrypted complete binary tree of depth D and returns a freshly generated key;
- `(tree-select key branches)` returns the subtree defined by the given branches;
- `(tree-mutate key subtree tx)` mutates the subtree defined by the given branches using the given transformation.

Additionally, we suppose in the following sections that the subtrees returned by this scheme are made of nodes that are either a vector containing the node data and its left and right children and sentinel nodes represented by the empty list are used to make processing the tree easier.

```
(define (new-node data l-node r-node) (vector data l-node r-node))
(define (sentinel) (list))
```

3.1 Setup procedure

The setup simply consists in calling the setup of the underlying (s, D) -binary-tree encryption scheme:

```
(define (ploc-setup) (tree-setup))
```

3.2 Insertion procedure

In PLOC as in most SSE, inserting requires the client to store the volumes of each label. This can be done using a hash-map:

```
(define volumes (hash-map))

(define (get-volume label)
  (hash-map-find volumes label 0))

(define (increment-volume label)
  (let ((volume (get-volume label)))
    (hash-map-bind! volumes label (+ volume 1))))
```

Each binding is stored wrapped inside a *datum* that holds both the indexed value, and a branch called *target*.

```
(define (datum value target)
  (vector value target))

(define (datum->value datum) (vector-ref datum 0))
(define (datum->target datum) (vector-ref datum 1))
```

The target is a pseudo-random number computed using the cryptographic hash H and the *key*.

```
(define (get-target key label volume)
  (H key label volume))
```

```
(define (make-datum key label value)
  (let* ((volume (get-volume label))
         (target (get-target key label volume)))
    (datum value target)))
```

Storing the target branch in each datum enables future modifications to *compact* the tree while maintaining the invariant (**I**) each datum is on its target branch, and the post-condition (**P**) no datum can be pushed further down.

The compaction is performed on a subtree of n branches that are generated using a procedure of no argument called a *scheduler*. Some scheduling strategies are better than others: if the same set of branches were to be constantly scheduled for compaction, the datums which target branches have a small intersection with this set of branches would accumulate near the root of the tree, eventually causing an overflow. We therefore choose a scheduling strategy that sweeps the tree branches in a regular and uniform way: the lower bits are all covered by at least one of the branches and each branch covers the remaining bits in a circular fashion. For example with $D = 3$ and $n = 2$, the two upper-most levels representing the two lowest bits are all covered. Then, at each call to the scheduler, these bits are prepended the bits generated by selecting the next integer, resulting in the subtrees depicted in Figure 1. The following `make-scheduler`

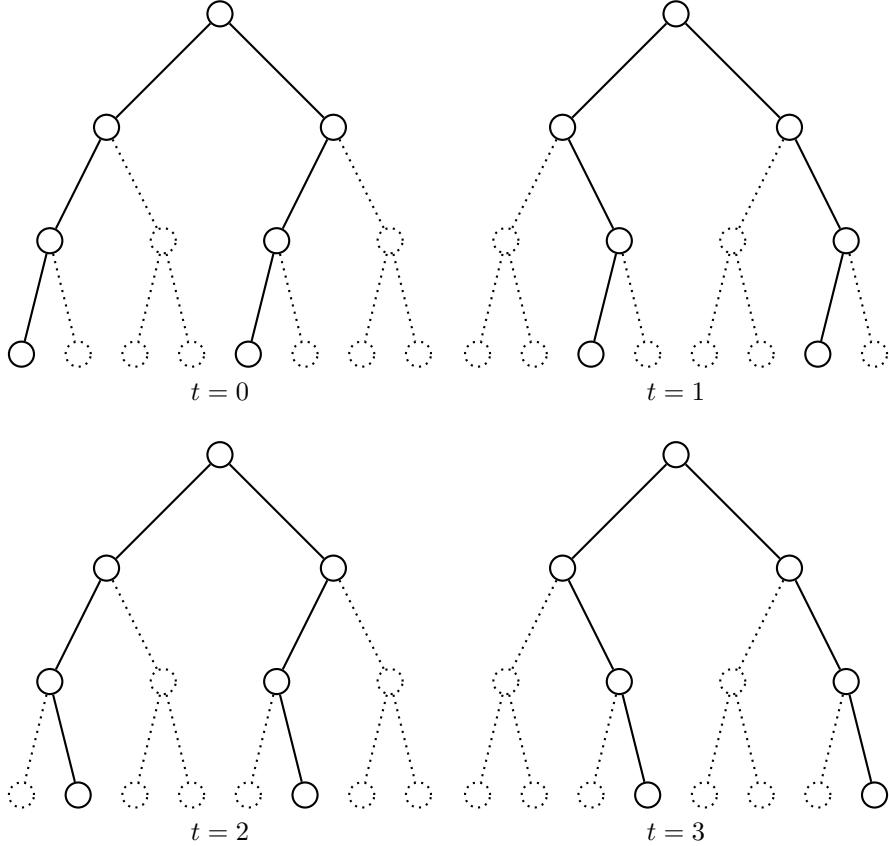


Figure 1: Scheduled subtrees for $D = 3$ and $n = 2$.

procedure generates a stateful scheduler drawing batches of n branches.

```
(define (make-scheduler n)
  (assert (integer? (lg n)))
  (let ((counter 0) (prefixes (iota n)))
    (λ ()
      (set! counter (+ counter 1))
      (map (λ (pfx) (+ pfx (* counter n)))
            prefixes))))
(define scheduler (make-scheduler n))
```

The compaction is a recursive procedure: it accumulates the data targeting each leaf on its way down, and stores it back to the tree on its way up, saturating the nodes with the c first remaining datums selected using the `take` procedure.

```
(define (take n xs)
  (let loop ((n n) (xs xs) (taken (list)))
    (if (or (zero? n) (null? xs))
        (values taken xs)
        (loop (- n 1) (cons (head xs) taken) (tail xs)))))
```

At each level, to tell apart data targeting the left and right subtrees, the procedure `(split-with pred? xs)` is used. It returns two lists: one containing values that match the given predicate and one containing those that don't.

```
(define (split-with pred? xs)
  (let loop ((xs xs) (lhs (list)) (rhs (list)))
    (match xs
      ((x . xs) (if (pred? x)
                     (loop xs (cons x lhs) rhs)
                     (loop xs lhs (cons x rhs))))
      (() (values lhs rhs)))))
```

The predicate `(go-left? depth branch)` returns whether the choice encoded at the given depth in the given branch is to select the left child.

```
(define (go-left? depth branch)
  (zero? (&& (<< 1 depth) branch)))

(define (compact node depth data)
  (define (left-datum? datum)
    (go-left? depth (datum->target datum)))
  (define (direct-data data)
    (split-with left-datum? data))
  (match node
    (#(node-data l-child r-child)
     (let*-values
       (((l-data r-data) (direct-data (append node-data data)))
        ((l-child l-data) (compact l-child (+ depth 1) l-data))
        ((r-child r-data) (compact r-child (+ depth 1) r-data))
        ((node-data rest) (take c (append l-data r-data))))
       (values (new-node node-data l-child r-child) rest)))
     (() (values (list) data))))
```

The preservation of the invariant **I**, can be proven by recurrence by showing that a datum is accumulated up to the last node in the intersection of its target branch with the fetched subtree. This proof relies on the correctness of the `left-datum?` and `split-with` procedures and the use of sentinel nodes that make processing the last node of the intersection like any other node. Once this is done, the rest of the proof follows trivially since the nodes in which a datum can be placed back are the nodes visited during the recursive calls which have been proven to belong to its target branch. Proving **P** is simple since the backtracking saturates the lower nodes first, until there is no datum left or the tree overflows.

In order to insert a new datum, PLOC simply calls the compaction procedure on a truly random subtree with this datum.

```
(define (ploc-insert key label value)
  (let* ((data (list (make-datum key label value)))
         (tx   (λ (tree)
                  (call-with-values (λ () (compact tree 0 data))
                    (λ (tree overflow)
                      (if (null? overflow)
                          tree
                          (error ploc-insert "compaction overflow"
                                 label value tree overflow)))))))
    (tree-mutate key (scheduler) tx)
    (increment-volume label)))
```

3.3 Search procedure

The search procedure is much simpler: it fetches the subtree generated by all branches possibly targeted by values indexed under the searched label,

```
(define (get-all-targets key label)
  (map (λ (v) (get-target key label v))
        (iota V)))
```

and retrieves the data targeting the same target branches:

```
(define (ploc-search key label)
  (let* ((targets (get-all-targets key label))
         (subtree (tree-select key targets))
         (values (tree-data subtree)))
    (reverse
      (foldl (λ (acc target)
                (match (multi-map-get values target)
                  (() acc)
                  ((value) (cons value acc))
                  (values (error ploc-search
                                 "more than one values found"
                                 target values))))
            (list)
            targets))))
```

where `tree-data` is the procedure using a depth-first search:

```
(define (tree-dfs f acc root)
  (let loop ((node root) (acc acc))
    (match node
      (#(data l-child r-child)
       (loop l-child
             (loop r-child
                   (f data acc))))
      (()) acc)))

```

to return a multi-map containing all values stored in the given tree, indexed by their target:

```
(define (tree-data tree)
  (define (index-datum mm datum)
    (match datum
      (#(value target)
       (multi-map-put mm target value))))
  (define (index-data mm data)
    (foldl index-datum mm data))
  (tree-dfs index-data (multi-map) tree)))

```

4 Security analysis

In this section, we assume the existence of an \mathcal{L}_τ -adaptively-secure (s, D) -binary-tree encryption scheme, where \mathcal{L}_τ is defined as follow:

- $\mathcal{L}_\tau(\text{Setup}) = \text{Setup}$;
- $\mathcal{L}_\tau(\text{Select}, \{b_i\}) = (\text{Select}, \{b_i\})$;
- $\mathcal{L}_\tau(\text{Mutate}, \{b_i\}, tx) = (\text{Mutate}, \{b_i\})$.

That is, these binary-tree-encryption-scheme operations only leak the branches defining the subtree they are working on. Such a scheme is actually trivial to design as it only requires nodes to be encrypted, and one possible implementation is described in appendix C.

Theorem 2. *Given an \mathcal{L}_τ -adaptively-secure (s, D) -binary-tree encryption scheme with public parameters pp_τ , PLOC is an \mathcal{L}_π -adaptively-secure (\mathbb{L}, \mathbb{V}) -multi-map encryption scheme with public parameters (pp_τ, c, n, B, V) , where $s = c \lg |\mathbb{V}|$, $D = \lg B$ and \mathcal{L}_π is defined as follows:*

- $\mathcal{L}_\pi(\text{Setup}) = \text{Setup}$
- $\mathcal{L}_\pi(\text{Search}, \ell) = (\text{Search}, \text{sp})$
- $\mathcal{L}_\pi(\text{Insert}, \ell, v) = \text{Insert}$

Proof. Given the domain \mathbb{L} and codomain \mathbb{V} , the public parameters (c, n, B, V) , and an \mathcal{L}_τ -adaptively-secure $(c \lg |\mathbb{V}|, \lg B)$ -binary-tree encryption scheme with public parameters pp_τ , there exists a simulator ϕ_{pp_τ} such that $\phi_{pp_\tau} \circ \mathcal{L}_\tau$ is indistinguishable from a legitimate tree-encryption-scheme client in the security game given in 2.4. In the following sections, we show how to build a simulator $\psi_{c,n,B,V}^{pp_\tau}$ such that $\psi_{c,n,B,V}^{pp_\tau} \circ \mathcal{L}_\pi$ is indistinguishable from a legitimate PLOC client in the same game. The full code of the simulator is given in section 4.5. \square

Note that in the implementation of the binary-tree encryption scheme given in appendix C, the only public parameter is the encryption overhead.

Corollary 1. *Given an \mathcal{L}_τ -adaptively-secure (s, D) -binary-tree encryption scheme with public parameters pp_τ , PLOC is an injection-secure (\mathbb{L}, \mathbb{V}) -multi-map encryption scheme with public parameters (pp_τ, c, n, B, V) , where $s = c \lg |\mathbb{V}|$, $B = \lg B$.*

Proof. Direct using the sufficient condition from [2]. \square

4.1 Setup

Recall the PLOC setup simply calls the tree setup, which only leaks the name of the operation:

```
(define (psi-setup) (phi 'setup))
```

Therefore, a PLOC setup only has trivial leakages:

$$\mathcal{L}_\pi(\text{Setup}) = \text{Setup}.$$

4.2 Search

In a nutshell, the only client-server interaction performed by a search consists in fetching the subtree targeted by the searched label. In particular, this subtree does not depend on the volume of this label but on V which is a public PLOC parameter. Then by security of the tree scheme, reading this subtree only leaks the branches that generate it, which in turn only depends on the search pattern which is therefore the only non-trivial leakage.

Recall the search operation works as follows:

```
(define (get-target key label volume)
  (H key label volume))

(define (get-all-targets key label)
  (map (λ (v) (get-target key label v))
    (iota V)))

(define (ploc-search key label)
  (let* ((targets (get-all-targets key label))
    (subtree (tree-select key targets))
    (values (tree-data subtree)))
  (reverse
    (foldl (λ (acc target)
      (match (multi-map-get values target)
        (() acc)
        ((value) (cons value acc))
        (values (error ploc-search
          "more than one values found"
          target values))))
      (list)
      targets))))
```

It can be transformed without adversarial advantage into a procedure ignoring the fetched tree and in which the calls to `get-target` and `get-all-targets` have been inlined:

```
@@ -1,21 +1,4 @@
-(define (get-target key label volume)
-  (H key label volume))
-
-(define (get-all-targets key label)
-  (map (λ (v) (get-target key label v))
-        (iota V)))
-
(define (ploc-search key label)
-  (let* ((targets (get-all-targets key label))
-         (subtree (tree-select key targets))
-         (values (tree-data subtree)))
-    (reverse
-     (foldl (λ (acc target)
-             (match (multi-map-get values target)
-                   (() acc)
-                   ((value) (cons value acc))
-                   (values (error ploc-search
-                                  "more than one values found"
-                                  target values))))
-            (list)
-            targets))))
+ (let* ((targets (map (λ (v) (H key label v)) (iota V)))
+        (subtree (tree-select key targets)))
+    (values)))
```

which gives:

```
(define (ploc-search key label)
  (let* ((targets (map (λ (v) (H key label v)) (iota V)))
         (subtree (tree-select key targets)))
    (values)))
```

In the oracle model, we can replace the keyed hash function with a negligible adversarial advantage by a call to a PRF:

```
@@ -1,4 +1,4 @@
(define (ploc-search key label)
-  (let* ((targets (map (λ (v) (H key label v)) (iota V))))
+  (let* ((targets (map (λ (v) (PRF label v)) (iota V))))
    (subtree (tree-select key targets)))
  (values)))
```

Then, by security of the PRF, the label can be replaced by any bijection $f : \mathbb{L} \rightarrow [1, |\mathbb{L}|]$, transforming the given `label` into its index `label-idx`:

```
@@ -1,4 +1,4 @@
-(define (ploc-search key label)
-  (let* ((targets (map (λ (v) (PRF label v)) (iota V))))
```

```

+(define (ploc-search key label-idx)
+  (let* ((targets (map (λ (v) (PRF label-idx v)) (iota V)))
         (subtree (tree-select key targets)))
      (values)))

```

which gives:

```

(define (ploc-search key label-idx)
  (let* ((targets (map (λ (v) (PRF label-idx v)) (iota V)))
         (subtree (tree-select key targets)))
      (values)))

```

Finally, by security of the binary-tree encryption scheme, the call to `tree-select` can be replaced by a call to $\phi_{pp_\tau} \circ \mathcal{L}_\tau$. Since $\mathcal{L}_\tau(\text{Fetch}, \text{targets}) = (\text{Fetch}, \text{targets})$, `(tree-select key targets)` is indistinguishable from `(phi 'select targets)`:

```

@@ -1,4 +1,4 @@
-(define (ploc-search key label-idx)
+(define (psi-search label-idx)
  (let* ((targets (map (λ (v) (PRF label-idx v)) (iota V)))
         -        (subtree (tree-select key targets)))
+        (subtree (phi 'select targets)))
      (values)))

```

which gives:

```

(define (psi-search label-idx)
  (let* ((targets (map (λ (v) (PRF label-idx v)) (iota V)))
         (subtree (phi 'select targets)))
      (values)))

```

Therefore, the only non-trivial leakage of a PLOC search is some indexing on the labels, which is the definition of leaking the search pattern:

$$\mathcal{L}_\pi(\text{Search}, \ell) = (\text{Search}, \text{sp}).$$

4.3 Insert

Recall the insertion works as follows:

```

(define (get-target key label volume)
  (H key label volume))

(define (make-datum key label value)
  (let* ((volume (get-volume label))
         (target (get-target key label volume)))
    (datum value target)))

(define (ploc-insert key label value)
  (let* ((data (list (make-datum key label value)))
         (tx   (λ (tree)
                  (call-with-values (λ () (compact tree 0 data))
                    (λ (tree overflow)
                      (if (null? overflow)

```

```

tree
(error ploc-insert "compaction overflow"
label value tree overflow)))))))
(tree-mutate key (scheduler) tx)
(increment-volume label)))

```

and suppose for now that compaction never overflows. The insertion can therefore be reduced without adversarial advantage to a procedure that ignores compaction error and does not modify the volumes:

```

@@ -10,10 +10,6 @@
(let ((data (list (make-datum key label value))))
  (tx  (λ (tree)
          (call-with-values (λ () (compact tree 0 data))
-           (λ (tree overflow)
-             (if (null? overflow)
-                 tree
-                 (error ploc-insert "compaction overflow"
-                       label value tree overflow)))))))
+           (λ (tree _) tree))))
  (tree-mutate key (scheduler) tx)
-  (increment-volume label)))
+  (values)))

```

By security of the tree encryption scheme, the call to `(tree-mutate key (scheduler) tx)` can be replaced with only negligible adversarial advantage by `(phi 'mutate (scheduler))`.

```

@@ -11,5 +11,5 @@
(tx  (λ (tree)
          (call-with-values (λ () (compact tree 0 data))
-           (λ (tree _) tree))))))
-  (tree-mutate key (scheduler) tx)
+  (phi 'mutate (scheduler))
  (values)))

```

which gives:

```

(define (get-target key label volume)
  (H key label volume))

(define (make-datum key label value)
  (let* ((volume (get-volume label))
         (target (get-target key label volume)))
    (datum value target)))

(define (ploc-insert key label value)
  (let ((data (list (make-datum key label value))))
    (tx  (λ (tree)
            (call-with-values (λ () (compact tree 0 data))
              (λ (tree _) tree))))))
  (phi 'mutate (scheduler))
  (values)))

```

Finally, removing unused variables can be performed without adversarial advantage:

```
@@ -1,15 +1,3 @@
-(define (get-target key label volume)
-  (H key label volume))
-
-(define (make-datum key label value)
-  (let* ((volume (get-volume label))
-         (target (get-target key label volume)))
-    (datum value target)))
-
-(define (ploc-insert key label value)
-  (let ((data (list (make-datum key label value))))
-    (tx  (λ (tree)
-            (call-with-values (λ () (compact tree 0 data))
-              (λ (tree _) tree))))))
-  (phi 'mutate (scheduler))
-  (values)))
+(define (psi-insert)
+  (phi 'mutate (scheduler))
+  (values))
```

which gives:

```
(define (psi-insert)
  (phi 'mutate (scheduler))
  (values))
```

Therefore a PLOC insertion only has trivial leakages if compaction does not overflow. The following section gives constraints on c to guarantee the absence of overflow with overwhelming probability in the security parameter λ .

4.4 Compaction overflow

We need to prove two different properties of the compaction process:

- a static property **S** guaranteeing that the tree does not overflow in its maximally compacted state. This is a classic requirement which most tree algorithms need to comply and has a well known solution in the balls-and-bins model (see for example [16]): it is enough to allocate $\frac{B}{N} + \frac{a \ln N}{\ln \ln N}$, $a > 1$ slots to each leaf in order to guarantee **S** with a probability in $1 - o(1)$, where N is the number of leaves in the tree, B the total number of values inserted (which we shall call bindings in our context). However, since the absence of overflow is critical for the rest of the security proof to hold, we will need to work out the proof of an upper bound with probability $1 - 2^{-\lambda}$.
- a dynamic property **D** guaranteeing that a congestion leading to an overflow at the root only happen with a probability negligible in λ in the worst case. To the best of our knowledge, this kind of probability is not standard.

First, recall that given a random variable X following a binomial law of parameter B and p , the Chernoff-Hoeffding theorem [9, 13] allows bounding the probability of the upper tail of that distribution:

$$\begin{aligned} P(X \geq k) &\leq \exp\left(-B \cdot D\left(\frac{k}{B} || p\right)\right) \\ &\leq \left[\left(\frac{k}{Bp}\right)^{\frac{k}{B}} \left(\frac{1-\frac{k}{B}}{1-p}\right)^{1-\frac{k}{B}}\right]^{-B} \\ &\leq \left(\frac{Bp}{k}\right)^k \left(\frac{1}{1-\frac{k}{B}}\right)^{B(1-\frac{k}{B})} \\ &\leq \left(\frac{Bp}{k}\right)^k e^{B\frac{k}{B}} \\ &\leq \left(\frac{eBp}{k}\right)^k \end{aligned}$$

In what follows, we will be interested in finding the smallest integer value of k such that $m \times P(X \geq k) \leq 2^{-\lambda}$ for some multiplicity m which value will depend on the proof. We therefore pose:

$$x(B, p, \lambda) (\ln(x(B, p, \lambda)) - \ln(Bp) - 1) = \lambda \ln(2) + \ln(m) \quad (1)$$

and call $k_{B,p}^{\lambda,m} = \lceil x(B, p, \lambda) \rceil - 1$, the smallest integer such that $P(X \leq k_{B,p}^{\lambda,m})$ holds with overwhelming probability.

4.4.1 Static bound

Since PLOC guarantees that the target leaves are selected using a PRF that is never called twice with the same inputs, we can consider the number of bindings targeting a leaf node to follow a binomial law of parameter $B = N$ and $p = N^{-1}$. Therefore we have:

$$P\left(\max_{1 \leq i \leq N} X_i \geq k\right) \leq N \times P(X_1 \geq k)$$

and maximum number of bindings received in a leaf is $k_{N,N-1}^{\lambda,N}$ with overwhelming probability in λ . Rather than allocating each leaf with this capacity, we rely on the fact that **S** is a static property that must only hold for the maximally compacted state of the tree. We can thus use slots from the parent nodes to store overflowing values. Let us count the number $\omega_c(c, d, b)$ of bindings overflowing from a node at level d knowing that b bindings have been inserted in its subtree and each node has a capacity c . In the worst case, all the bindings are stored in a contiguous region of the tree. Calling $k_{b,d}^\lambda = k_{b,2^{-d}}^{\lambda,2^d}$, we can bound $\omega_c^\lambda(b, d)$ by the overflow of its left subtree in which at most $k_{B,d+1}^\lambda$ bindings have been inserted, plus the overflow of its right subtree in which the remaining bindings have been inserted, minus its capacity c :

$$\omega_c^\lambda(b, d) = \omega_c^\lambda(\min(b, k_{B,d+1}^\lambda), d+1) + \omega_c^\lambda(b - \min(b, k_{B,d+1}^\lambda), d+1) - c,$$

with the border condition: $\omega_c^\lambda(b, \lg N) = b - c$. Hence, guaranteeing **S** can be done by choosing a value of c such that $\omega_c^\lambda(B, 0) = 0$, which is easy using a numerical application (cf tables 2 and 3 below).

4.4.2 Dynamic bound

While PLOC has been defined to be independent from the scheduler used to select the subtree to fetch upon each insertion, the proof of **D** given here heavily relies on the use of the uniform scheduler presented in section 3.2. Recall that this scheduler guarantees that at each depth d , a node is alternatively selected with its left and right child. The compaction period of each node at a depth d is thus $T(n, d) = \frac{2^d}{n}$, where n is the number of branches selected by the scheduler and must be a power of two. The second and more important consequence is that in the absence of congestion, the lifetime of a datum in any given node cannot be more than one compaction cycle. Indeed, if a datum is placed in a given node at a given compaction, it means the next node on its target path is the child of this node that hasn't been selected for compaction. At the next compaction cycle, that child will be selected and this datum pushed along its target path at least one level down.

With this in mind, consider for once that nodes have an unbounded capacity. For the same reason as in the static proof, we can consider the number of bindings held by a node at depth $d < \lg(N)$ to follow a binomial law of parameter $T(n, d)$ and $p(d+1)$, where $p(d) = 2^{-d}$. Indeed, $T(n, d)$ bindings have been inserted in the tree since the last compaction, and a binding targets that node if and only if (unbounded node capacity) the next node in its target path is that node's child that hasn't been selected for compaction. During a compaction, we are modifying at most $m = n \lg N$ nodes and we therefore need to use a node capacity $c(\lambda, n, N)$ guaranteeing the absence of overflow for that many nodes at each level: for example, $\max_{0 \leq d < \lg N} k_{T(n, d), p(d+1)}^{\lambda, n \lg N}$. However, nodes do have a bounded capacity and we still need to prove that no node ever overflows even in face of the worst insertion sequence.

We define an overflowing path to be an overflowing leaf or an overflowing node which has at least one overflowing path among its children, and prove by recurrence the property (\mathcal{P}) : in a tree of N leaves, a node of capacity $c(N, \lambda)$ that does not belong to an overflowing path overflows with a probability negligible in λ .

- (\mathcal{P}) trivially holds for all leaves.
- Let $d > 1$ be a depth at which \mathcal{P} holds for all nodes. Upon compaction of a node at depth $d - 1$, that node receives a set of new bindings that are merged with its current bindings. Without loss of generality, let's call left child the one that has been selected for compaction. This child receives all the bindings received by its parents that target it.
 - If it overflows, since (\mathcal{P}) holds at depth d , this child belongs to an overflowing path and an overflow of its parent would extend this path.
 - If it not overflow, the set of bindings stored in its parent is exactly the set of all old and new bindings targeting the right child. On the one hand, a non-empty set of old bindings means the right child overflowed during the previous compaction. Since (\mathcal{P}) holds at the depth d to which this child belongs, it also belongs to an overflowing path that would be extended by an overflow of its parent. On the other hand, if the set of old bindings is empty, the node capacity $c(N, \lambda)$ guarantees the absence of overflow with overwhelming probability.

Therefore, a tree overflow can only happen if the root is the tip of an overflowing path to some leaf. In particular, if the leaves cannot overflow, the tree cannot overflow. We could choose the leaf capacity to be $k_{N,N-1}^{\lambda,N}$ to guarantee **D**. It would however unnecessarily increase the tree size (we achieved to prove **S** for smaller leaf sizes) and we instead prove that, given nodes of capacity c^* , there exists d^* such that no subtree of depth d^* which leaves are leaves of the main tree can overflow. In particular, both **S** and **D** need to hold for those subtrees. c^* must therefore be the maximum of the capacities required for each property to hold: $c^* = \max(c_s^*, c_d^*)$.

In what follows, we prove that there exists a depth d^* such that the number $2^i B^*$ of bindings arriving in any subtree of depth d^* during the time required to compact i levels is smaller than ic^* , where $B^* = T(n, d_{\max} - d^*)$, $p^* = p(d_{\max} - d^*)$ and $c^* = k_{B^*, p^*}^{\lambda, m}$. By definition, we have for each $k > k_{B^*, p^*}^{\lambda, m}$:

$$m \left(\frac{e\mu}{k} \right)^k \leq 2^{-\lambda}$$

and considering the equality case:

$$k(\ln(k) - \ln(\mu) - 1) = \lambda \ln(2) + \lg(m).$$

Supposing k is a continuous function in i and $\mu = 2^i B^*$, and taking the partial derivative with respect to i of both sides we have:

$$\begin{aligned} & \frac{\partial k}{\partial i} (\ln(k) - \ln(\mu) - 1) + k \left(\frac{1}{k} \frac{\partial k}{\partial i} - \frac{1}{\mu} \frac{\partial \mu}{\partial i} \right) = 0 \\ \iff & \frac{\partial k}{\partial i} \left(\frac{\lambda \ln(2) + \ln(m)}{k} \right) + \frac{\partial k}{\partial i} - \frac{k}{\mu} \frac{\partial \mu}{\partial i} = 0 \\ \iff & \frac{\partial k}{\partial i} \left(\frac{\lambda \ln(2) + \ln(m)}{k} + 1 \right) - \frac{k}{\mu} \frac{\partial \mu}{\partial i} = 0 \\ \iff & \frac{\partial k}{\partial i} = \frac{\partial \mu}{\partial i} \left(\frac{\frac{k}{\mu}}{\frac{\lambda \ln(2) + \ln(m)}{k} + 1} \right) \\ \iff & \frac{\partial k}{\partial i} = \frac{\partial \mu}{\partial i} \left(\frac{k^2}{\mu(k + \lambda \ln(2) + \ln(m))} \right) \\ \iff & \frac{\partial k}{\partial i} = \frac{k^2 \ln(2)}{k + \lambda \ln(2) + \ln(m)} \end{aligned}$$

Therefore,

$$\begin{aligned} & \frac{\partial k}{\partial i} \leq c^* \\ \iff & k^2 - \frac{c^*}{\ln(2)}(k + \lambda \ln(2) + \ln(m)) \leq 0 \\ \iff & k \leq \frac{c^*}{2 \ln(2)} \left(1 + \sqrt{1 + \frac{4 \ln(2)}{c^*}(\lambda \ln(2) + \ln(m))} \right) \end{aligned}$$

Noting k^* the biggest integer k verifying that inequality, and $d^* = \frac{k^*}{c^*}$, we have for each $i \leq d^*$, $k(i) \leq ic^*$ which proves **D** in any subtree of depth d^* . Proving **S** in that subtree can be done by modifying the proof given in section 4.4.1 in order to extract the first overflowing depth and ensuring it is bigger than $d_{\max} - d^*$. The following table 2 and 3 give some values of the node capacity c^* for which both **S** and **D** hold for some values of λ , n and N :

4.5 PLOC simulator

Given phi, the simulator of the tree scheme, we can now build the simulator of the client-side PLOC scheme. We have proven in the previous sections that this

Table 2: $c^*(\lambda, n, N)$ in function of n and D for $\lambda = 128$.

n	D		
	16	32	40
16	19	19	19
32	17	17	17
64	16	16	16
128	14	14	14
256	13	13	13

Table 3: $c^*(\lambda, n, N)$ in function of n and D for $\lambda = 64$.

n	D		
	16	32	40
16	11	11	11
32	10	10	10
64	9	9	9
128	8	8	8
256	7	8	8

simulator is indistinguishable from a legitimate client. Since it only requires the name of the operation to perform and some indexing function of the label upon search, this proves that PLOC is injection-secure.

```
(define-module (sse ploc simulator)
  #:use-module (sse utils)
  #:use-module (sse ploc scheduler))

(define (make-psi pp-tree c n B V value-size PRF make-phi)

  (define phi (make-phi pp-tree (* c value-size) (lg B)))

  (define scheduler (make-scheduler n))

  (define (psi-setup) (phi 'setup))

  (define (psi-search label-idx)
    (let* ((targets (map (λ (v) (PRF label-idx v)) (iota V)))
           (subtree (phi 'select targets)))
        (values)))

  (define (psi-insert)
    (phi 'mutate (scheduler))
    (values))

  (λ (op . args)
    (match op
      ('setup (apply psi-setup args))
      ('search (apply psi-search args))
      ('insert (apply psi-insert args))))
```

5 Performance analysis and variations on PLOC

Theorem 3. *When instantiated using an $(c \lg |\mathbb{V}|, B)$ -tree encryption scheme with $O(|\{b_i\}| \times c \lg |\mathbb{V}| \lg B)$ Fetch and Merge operations, PLOC is an (\mathbb{L}, \mathbb{V}) -multi-map encryption scheme with two-round $O((c \lg |\mathbb{V}| + C_1)n \lg B)$ Insert operations and one-round $O((c \lg |\mathbb{V}| + C_2)V \lg B)$ Search operations.*

Proof. The number of rounds can be enumerated from the code of the reference implementation described in section 3. In the case of an Insert operation, a subtree of n leaves is successively fetched, modified locally, then merged back to the main tree. Since the tree encryption scheme is instantiated with $s = c \lg |\mathbb{V}|$ and $D = \lg B$, the complexity of both tree operations is $O(c \lg |\mathbb{V}| \times n \lg B)$. The compaction of the tree being linear in its size, its complexity is in $O(n \lg B)$. In the case of a Search operation, a single tree Fetch is performed on a subtree of V leaves, which gives the desired complexity. \square

Theorem 3 guarantees logarithmic performances in B for PLOC, although with a non-negligible constant. In practice, the performances of the reference implementation given here – which is optimized for clarity rather than speed – are good: on an Intel(R) Core(TM) i5-6200U CPU @ 2.30GHz, we can measure the following timings using GNU Guile v3.10:

Table 4: Operation performances (Search Insert) in function of n and B for $V = \sqrt{B}$ and c given in Table 2.

$n \setminus B$	2^{10}		2^{16}		2^{20}	
16	2.1 msec	1.8 msec	26 msec	3.1 msec	0.17 sec	3.8 msec
64	2.1 msec	5.4 msec	25 msec	11 msec	0.12 sec	13 msec
256	2.0 msec	18 msec	25 msec	34 msec	0.12 sec	45 msec

Note that increasing n allows decreasing c , but it does not impact the Search timings much since nodes are mostly padding and the PLOC-operation constant is much bigger than the tree-operation constant ($C_2 \gg 1$). The Search timings are mostly impacted by V which is linked to B in the present benchmark.

5.1 Fully-dynamic multi-map encryption scheme

The scheme described in section 3 only supports Insert and Search queries. However, the possibility to delete bindings is usually needed by concrete use-cases. It is actually possible to implement a fully-dynamic multi-map encryption scheme on top of PLOC with the same security and performance characteristics using a classic technique known as *lazy deletions* that we find more appropriate to call *journaling*: instead of inserting bindings in a multi-map, we use this multi-map to store per-label log entries containing arbitrary operations. This technique allows implementing any associative abstract data type.

Definition 5.1 (Mapping). An associative abstract data type (or mapping) is an abstract data type:

- supporting a valid empty state;

- supporting an operation `Search` parameterized by an argument type called labeling domain, noted \mathbb{L} , which sole purpose is to retrieve the value(s) currently indexed under a given label ℓ ;
- of which all other operations are parameterized by a composite argument type $\mathbb{A} = \mathbb{L} \times \mathbb{A}'$ and the result type $\mathbb{R} = \{\perp\}$, and have for sole purpose to mutate the value(s) indexed under some label ℓ .

Remark. A multi-map is a mapping for which $\mathbb{A}' = \mathbb{V}$ and the parametric result type for the search operation is $\mathbb{V}^* = (\perp \cup \mathbb{V}) \times \mathbb{V}^*$.

Logging a mapping operation is as simple as inserting the operation name and its rest arguments (in \mathbb{A}') under the targeted label:

```
(define (encrypted-mapping-mutate key op label args)
  (ploc-insert key label (list op args)))
```

Note that mapping mutations – once specialized with their arguments – consume and return a mapping data structure and therefore form a monoid which neutral element is the identity. This property allows us to reduce the log entries to a unique mapping transformation that can then be used to produce the current mapping state by applying it to the empty state:

```
(define (specialize entry)
  (let ((op (head entry))
        (args (tail entry)))
    (op label args mapping)))

(define (encrypted-mapping-search key label)
  (let* ((log (ploc-search key label))
         (ops (map specialize log))
         (tx (reduce compose identity ops)))
    (mapping-search label (tx empty-mapping))))
```

Corollary 2. *There exists an injection-secure fully-dynamic multi-map encryption scheme with the same performance characteristics as PLOC.*

Proof. Direct using the journaling transformation proposed above. □

Remark. This transformation allows implementing an injection-secure fully-dynamic multi-map encryption scheme, but does not allow reclaiming the space used by the deleted bindings. To that end, we need to add a special garbage-collection primitive to PLOC that reads the subtree targeted by a given label, extracts the values bound to this label, filters out the deleted ones, compacts the subtree with the remaining bindings, merges the subtree back to the main tree and returns the remaining values to the client. Like the search operation, garbage collection only leaks the permuted label index and can therefore be executed few such operations instead of the search, in order to amortize the cost of its additional round-trip.

5.2 Stateless insertions

Corollary 3. *There exists an injection-secure fully-dynamic multi-map encryption scheme with stateless client, four-round $O((c \lg |\mathbb{V}| + C_1)n \lg B + C'_1 \lg L)$*

Insert *operations*, where L is an upper-bound on the number of labels, and $O((c \lg |\mathbb{V}| + C_2)V \lg B)$ non-interactive Search operations.

Informal proof: It is possible to get a stateless ORAM client by storing the padded stash server-side, for example using PathORAM [17]. When using it to store the label volumes, the costs a querying a label is $O(\lg L)$. The protocol for performing a mutation would be:

1. retrieve the stash from the server: $O(\lg L)$;
2. read the ORAM in search for the volume of the mutated label: $O(\lg L)$;
3. perform the PLOC mutation $O((c \lg |\mathbb{V}| + C_1)n \lg B)$;
4. write back the updated stash and ORAM data structure: $O(\log L)$.

Since both the ORAM and the PLOC mutation have no non-trivial leakages, accomplishing these four steps only has trivial leakage, which allows preserving PLOC's injection security.

6 Conclusion and future work

In this article, we have presented PLOC, a new SSE scheme, and proved that it is injection-secure. Doing so, we achieved state-of-the-art security while significantly improving upon the only other injection-secure SSE scheme known to date. Our proof of injection security relies on a new way to decorrelate mutations from their effect: we rely on the root being a member of all branches to push store values in the path to their target leaf without having to fetch branches derived from those values nor their indexing label. Central to this proof is the fact that compaction does not overflow. To the best of our knowledge, there has been no study of such mechanism and we therefore give the first bound on the node capacity allowing to prevent overflow with overwhelming probability. In a future work, we would like to see that bound tightened in order to improve the concrete performance of our scheme.

The stateless nature of PLOC's search operations makes it trivially amenable to a single-writer, multi-reader setting without loosing any security. We believe it also amenable to the fully-concurrent setting, which would be a break-through as the only fully concurrent scheme known to date does not even guarantee forward security in all cases. In order to achieve this, the client-side counters (only required for insertion) needs to be synchronized or removed. One way to remove it would be to prove that there exists a node capacity such that inserting values targeting a random branch from the target set of their indexing label does not cause the tree to overflow.

Finally, we believe that PLOC has the potential to bridge the gap between security and real-world applications with stringent performance requirements. We therefore hope to see used in many applications.

7 Acknowledgments

The authors would like to thank David Pointcheval for his invaluable feedback and the numerous and insightful shared discussions, Cosmian for funding long-term research and the Cosmian security team for their support.

References

- [1] Archita Agarwal, Seny Kamara, and Tarik Moataz. “Concurrent Encrypted Multimaps”. In: 2024, pp. 169–201. DOI: 10.1007/978-981-96-0894-2_6.
- [2] Ghous Amjad, Seny Kamara, and Tarik Moataz. “Injection-Secure Structured and Searchable Symmetric Encryption”. In: 2023, pp. 232–262. DOI: 10.1007/978-981-99-8736-8_8.
- [3] Ghous Amjad et al. “Dynamic Volume-Hiding Encrypted Multi-Maps with Applications to Searchable Encryption”. In: 2023.1 (Jan. 2023), pp. 417–436. DOI: 10.56553/popets-2023-0025.
- [4] Raphael Bost. “ oφoς : Forward Secure Searchable Encryption”. In: *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*. CCS ’16. Vienna, Austria: Association for Computing Machinery, 2016, 1143–1154. ISBN: 9781450341394. DOI: 10.1145/2976749.2978303. URL: <https://doi.org/10.1145/2976749.2978303>.
- [5] Raphael Bost and Pierre-Alain Fouque. “Security-Efficiency Tradeoffs in Searchable Encryption”. In: 2019.4 (Oct. 2019), pp. 132–151. DOI: 10.2478/popets-2019-0062.
- [6] Raphaël Bost, Brice Minaud, and Olga Ohrimenko. “Forward and Backward Private Searchable Encryption from Constrained Cryptographic Primitives”. In: *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*. CCS ’17. Dallas, Texas, USA: Association for Computing Machinery, 2017, 1465–1482. ISBN: 9781450349468. DOI: 10.1145/3133956.3133980. URL: <https://doi.org/10.1145/3133956.3133980>.
- [7] Théophile Brézot and Chloé Hébant. *Findex: A Concurrent and Database-Independent Searchable Encryption Scheme*. Cryptology ePrint Archive, Report 2024/1541. 2024. URL: <https://eprint.iacr.org/2024/1541>.
- [8] Melissa Chase and Seny Kamara. “Structured Encryption and Controlled Disclosure”. In: *Advances in Cryptology - ASIACRYPT 2010*. Vol. 6477. Lecture Notes in Computer Science. Springer Verlag, 2010, pp. 577–594.
- [9] Chernoff. “A Measure of Asymptotic Efficiency for Tests of a Hypothesis Based on the sum of Observations”. In: 23 (1952). DOI: 10.1214/aoms/1177729330.
- [10] Reza Curtmola et al. “Searchable symmetric encryption: improved definitions and efficient constructions”. In: *Proceedings of the 13th ACM Conference on Computer and Communications Security*. CCS ’06. Alexandria, Virginia, USA: Association for Computing Machinery, 2006, 79–88. ISBN: 1595935185. DOI: 10.1145/1180405.1180417.
- [11] Ioannis Demertzis et al. “Dynamic Searchable Encryption with Small Client Storage”. In: 2020. DOI: 10.14722/ndss.2020.24423.
- [12] R. Kent Dybvig. *The Scheme Programming Language, 4th Edition*. 4th. The MIT Press, 2009. ISBN: 026251298X.
- [13] Wassily Hoeffding. “Probability inequalities for sum of bounded random variables”. In: 1963. URL: <https://api.semanticscholar.org/CorpusID:121341745>.

- [14] Evgenios M. Kornaropoulos et al. “Leakage Inversion: Towards Quantifying Privacy in Searchable Encryption”. In: *Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security*. CCS ’22. Los Angeles, CA, USA: Association for Computing Machinery, 2022, 1829–1842. ISBN: 9781450394505. DOI: 10.1145/3548606.3560593. URL: <https://doi.org/10.1145/3548606.3560593>.
- [15] Sarvar Patel et al. “Mitigating Leakage in Secure Cloud-Hosted Data Structures: Volume-Hiding for Multi-Maps via Hashing”. In: *Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security*. CCS ’19. London, United Kingdom: Association for Computing Machinery, 2019, 79–93. ISBN: 9781450367479. DOI: 10.1145/3319535.3354213. URL: <https://doi.org/10.1145/3319535.3354213>.
- [16] Martin Raab and Angelika Steger. ““Balls into Bins” — A Simple and Tight Analysis”. In: *Randomization and Approximation Techniques in Computer Science*. Ed. by Michael Luby, José D. P. Rolim, and Maria Serna. Berlin, Heidelberg: Springer Berlin Heidelberg, 1998, pp. 159–170. ISBN: 978-3-540-49543-7.
- [17] Emil Stefanov et al. “Path ORAM: An Extremely Simple Oblivious RAM Protocol”. In: *J. ACM* 65.4 (Apr. 2018). ISSN: 0004-5411. DOI: 10.1145/3177872. URL: <https://doi.org/10.1145/3177872>.

A Reflection on the making of this article

Besides proposing a novel and powerful SSE scheme, this article also innovates on the way it has been produced, which we discuss in the following sections. We hope this can inspire other cryptographers and that the practices adopted here become more widespread in the future.

A.1 Literate programming

Literate programming is an implementation technique pioneered by Donald Knuth. It inverts the usual paradigm by inserting fragments of code in a text document instead of inserting fragments of documentation in the source files. The aim is to expose the code in an order that better fits the prose. Code fragments are then extracted and *tangled* back together by a dedicated software in order to produce the source files required by the compiler.

This article has been written in Org – an expressive markup language – with Emacs’ Org Mode which provides a way to both tangle code fragments and to export a document containing source code to different back-ends like L^AT_EX or HTML. That way the `.org` file is self-sufficient, which helps to make the article reproducible.

A.2 Using `pseudo` code

Rather than exposing our algorithm using some pseudo-code flavor, we chose to directly expose its implementation in a programming language to inherit from a well-defined – here standardized – semantics. Three things helped to achieve that: literate programming to make it possible, Emacs’ Org Mode which made

it easy and the Scheme programming language which allows expressing complex algorithms in few lines.

A.3 A path to formal verification

Apart from unambiguousness, using a programming language to present the algorithm and deriving a simulator has direct, very desirable side-effects:

- it allows compiling and testing the very algorithm described in the paper to verify it is valid, exhaustive and to test it performs as intended;
- it provides a reference implementation which is explained in great details;
- it simplifies the review process of our proofs since we used code diffs to highlight the changes. These diffs are automatically generated, allowing readers to fully trust them: there is no hidden, unexplained additional change in-between code transformations;
- it allows running the simulator given in section 4.5 against the reference implementation to test their indistinguishability.

This approach opens a way in-between publishing without verified source code and publishing automated verification of a source code, which often makes up for a paper in itself: our implementation is formally verified, but this verification is not automated and relies on careful reviews of our proof and the indistinguishability tests of our reference implementation.

B A Scheme primer

Scheme is a functionally-oriented programming language which simple syntax and expressiveness make it a target of choice to express algorithms. In this section, we provide the basics required for readers that are not familiar with this language to understand the PLOC implementation described in this article. For a complete reference, see [12].

B.1 Functions

Scheme implements the untyped lambda calculus, which defines the following operations:

- $(\lambda x.y)$ (abstraction) is the function with formal parameter x and body y ;
- $(M N)$ (application) evaluates the function M with the value N as parameter;

Similarly, a Scheme function is declared using the λ operator:

```
(λ (arg ...) body ...)
```

in which the formal arguments arg are bound to values upon application using the syntax:

```
((λ (arg ...) body ...) val ...)
```

For example, $((\lambda (x y) (+ x y)) 1 2)$ evaluates to 3. The main difference here between Scheme and the lambda calculus is that Scheme supports the definition of functions of multiple arguments while the untyped lambda calculus does not and that Scheme functions support returning multiple values using **values**.

B.2 Bindings

Departing from the lambda calculus, Scheme supports binding names to values. For historical reasons, there are many bindings operators:

- `(define name value)` binds `name` to `value` in the enclosing scope;
- `(define-values (name ...) values)` same as `define` but is used to bind values produced by `values`.
- `(let ((name value) ...) body ...)` generates a new scope in which the given names are bound to the given values, and in which the body is evaluated;
- `(let* ((name value) ...) body ...)` same as `let` but each binding in the scope of the following one;
- `(let-values (((name ...) values) ...) body ...)` same as `let` but is used to bind values produced by `values`;
- `(let*-values (((name ...) values) ...) body ...)` same as `let*` but is used to bind values produced by `values`.

For example, a function can be named using `define`:

```
(define add (\lambda (x y) (+ x y)))
```

which has an equivalent syntactic sugar:

```
(define (add x y) (+ x y))
```

And any symbol can be rebound using `set!`:

```
(set! add (\lambda (x y z) (+ x y z)))
```

A binding syntax of particular interest call named `let`, allows combining the definition of a recursive function and its initial call:

```
(let name ((arg init) ...)
  body ...)
```

is equivalent to:

```
(letrec ((name (\lambda (arg ...) body ...)))
  (name init ...))
```

B.3 Types

Although modeled after the untyped lambda calculus, type information is available at runtime. There are mainly four data-types in use in this article: boolean values, numbers, lists and vectors:

- boolean values are either `#t` (true) or `#f` (false). They support the usual boolean operators `and` and `or` and are used for control flow. For example, `(if boolean body1 body2)` evaluates to `body1` if `boolean` is `#t` and `body2` otherwise.
- numbers can be arbitrary large and support all the usual operations like `+` and `modulo`;
- lists are generated using `(list value ...)` and support dynamic resizing via prepending with `(cons value lst)` which returns a list starting with `value` and ending with the same elements as `lst`, and appending via `(append lst ...)` which returns a list starting with the elements in `lst` and ending with the elements of the subsequent lists.
- vectors are generated using `(vector value ...)` and do not support dynamic resizing.

List elements are accessed using `(list-ref lst index)` and vector elements are indexed using `(vector-ref vec index)`.

Scheme supports pattern matching using `(match value (pattern body) ...)` in which the body evaluated is the one associated to the first matching pattern. These patterns can be:

- any boolean or number, matching this value;
- `()` matching a list of no element;
- `(head . tail)` matching a list of at least one element with first element `head` and remaining element list `tail`, in which case `body` is evaluated with `head` and `tail` in scope;
- `#{element ...}` matching a vector containing the given number of elements, in which case `body` is evaluated with these elements in scope.

For example, the following piece of code taken from the PLOC implementation defines a function named `split-with` that takes two arguments: a predicate `pred?` and a list `xs`. It deconstructs `xs` until it is reduced to the empty list, successively prepending each deconstructed element from `xs` to `lhs` if this element matches the predicate `pred?` or `rhs` otherwise. Finally, it returns both `lhs` and `rhs`.

```
(define (split-with pred? xs)
  (let loop ((xs xs) (lhs (list)) (rhs (list)))
    (match xs
      ((x . xs) (if (pred? x)
                     (loop xs (cons x lhs) rhs)
                     (loop xs lhs (cons x rhs))))
      (()       (values lhs rhs))))
```

B.4 Classical procedures

The PLOC implementation also makes use of the classical procedures:

- (`iota n`) which returns the list `(0 1 ... (- n 1))`;
- (`foldl f a xs`) which accumulates each element of `xs` using `(f a x)` then returns `a`;
- (`map f xs`) which returns the list of all values in `xs` transformed using `(f x)`;
- (`filter pred? xs`) which returns the list of all values in `xs` matching the predicate function `pred?`.

C Tree Encryption Scheme

PLOC relies on an (s, D) -binary-tree encryption scheme, defined in section 3. The present section describes an implementation of this scheme for completeness only.

C.1 Tree serialization

Following the design choices discussed in [7], we rely on a server abstraction exposing a linear 128-bit memory space with the following operations:

- (`memory-setup`) initializes a new server-side abstraction of a 128-bit linear memory space;
- (`memory-read addresses`) atomically reads from given addresses and returns the list of words bound to these addresses, or the special value `free` in case no word was found;
- (`memory-write bindings`) atomically writes the given bindings to memory.

In order to fit a tree into a linear memory, we need to associate each node to an address that is computed using its target and depth as follows:

```
(define (make-address depth target)
  (G depth (and target (- (<< 1 depth) 1))))
```

where `G` is a hash function which is only required to have a good-enough collision resistance.

C.2 Setup

The `tree-setup` simply consists in setting-up a new server-side memory, and returning a fresh AE key.

```
(define (tree-setup)
  (memory-setup)
  (AE-keygen))
```

Not initializing the tree avoids a $O(N)$ setup without leaking any additional information in the persistent-adversary model. Note that we could pre-allocate enough server-side memory up-front since the size of the tree is already known.

C.3 Select

To return the subtree generated by the given target branches, we first need to derive all the addresses of its nodes, which can be performed using the following `subtree-addresses` procedure:

```
(define (subtree-addresses branches)
  (let loop ((depth 0) (branches branches))
    (if (or (null? branches) (< max-depth depth))
        (list)
        (let-values
            (((lhs rhs) (split-with (λ (b) (go-left? depth b))
                                      branches)))
             (cons (make-address depth (car branches))
                   (append (loop (+ depth 1) lhs)
                           (loop (+ depth 1) rhs)))))))
```

Reading data from a word is performed by decrypting this word and interpreting the resulting bytes as a serialized data. Since the tree was not initialized in-memory upon setup, some addresses might be free, which must be taken into account:

```
(define (word->data AE word)
  (if (null? word)
      (list)
      (bytes->data (AE-decrypt AE word))))
```

Finally, the tree can be built from the words read by relying on the fact that `memory-read` preserves the address order: we can perform the exact *reverse* tree traversal as in `subtree-addresses` (in which the right recursive call is performed first while it was performed last in `subtree-address`) since the list in which the addresses were accumulated is a stack.

```
(define (words->tree AE branches words)
  (let loop ((depth 0) (branches branches) (words words))
    (if (or (null? branches) (< max-depth depth))
        (values (sentinel) words)
        (let*-values
            (((lhs      rhs) (split-with
                               (λ (b) (go-left? depth b))
                               branches))
             ((word      words) (values (car words) (cdr words)))
              ((r-child words) (loop (+ depth 1) rhs words)))
              ((l-child words) (loop (+ depth 1) lhs words)))
             (values (new-node (word->data AE word)
                               l-child r-child)
                     words))))))
```

We can now write the complete `tree-select` procedure:

```
(define (tree-select key branches)
  (let* ((addresses (subtree-addresses branches))
         (words      (memory-read addresses)))
    (words->tree (AE-init key) branches words)))
```

C.4 Mutate

Mutating the tree consists in fetching the subtree defined by the given branches much like the `select` procedure, generating the new subtree using the given transformation and writing it back to memory. To do that, we first need a procedure to convert a tree into bindings:

```
(define (tree->bindings AE root)
  (let loop ((depth 0) (branch 0) (node root))
    (match node
      (#(data l-child r-child)
       (let ((address (make-address depth branch)))
         (word (AE-encrypt AE (data->bytes data))))
       (cons (cons address word)
             (append (loop (+ depth 1) (+ (<< branch 1) 0) l-child)
                     (loop (+ depth 1) (+ (<< branch 1) 1) r-child))))))
      () (list)))))

(define (tree-mutate key branches transformation)
  (let* ((addresses (subtree-addresses branches))
         (words (memory-read addresses))
         (AE (AE-init key))
         (old-tree (words->tree AE branches words))
         (new-tree (transformation old-tree)))
    (memory-write (tree->bindings AE new-tree))))
```

D Reference implementation

D.1 Useful functions

We first define some additional wrappers, utility functions and data-structures.

```
(library (sse utils)
  (export λ lg << >> && u8? u64? iota foldl split-with take match
          hash-map hash-map-find hash-map-bind!
          multi-map multi-map-get multi-map-put!
          new-node sentinel go-left? tree-dfs)
  (import (rnrs base)
          (rnrs control)
          (rnrs lists)
          (rnrs hashtables)
          (rnrs arithmetic bitwise)
          (only (ice-9 match) match)
          (only (srfi :1) iota))

  (define-syntax λ
    (syntax-rules ()
      ((_ expr ...) (lambda expr ...)))))

  (define (u8? n)
    (and (integer? n)
```

```

(<= 0 n) (< n (<< 1 8)))))

(define (u64? n)
  (and (integer? n)
    (<= 0 n) (< n (<< 1 64)))))

(define (lg n) (/ (log n) (log 2)))

(define (<< n e) (bitwise-arithmetic-shift-left n e))
(define (>> n e) (bitwise-arithmetic-shift-right n e))
(define (&& n m) (bitwise-and n m))

(define head car)
(define tail cdr)

(define foldl fold-left)

(define (take n lst)
  (let loop ((n n) (pxf (list)) (lst lst))
    (if (or (zero? n) (null? lst))
        (values pxf lst)
        (loop (- n 1) (cons (head lst) pxf) (tail lst)))))

(define (split-with pred? xs)
  (let loop ((xs xs) (lhs (list)) (rhs (list)))
    (if (pair? xs)
        (if (pred? (head xs))
            (loop (tail xs) (cons (head xs) lhs) rhs)
            (loop (tail xs) lhs (cons (head xs) rhs)))
        (values lhs rhs)))))

(define (hash-map) (make hashtable equal hash equal?))
(define (hash-map-find hmap k d) (hashtable-ref hmap k d))
(define (hash-map-bind! hmap k v) (begin (hashtable-set! hmap k v) hmap))

(define (multi-map) (hash-map))
(define (multi-map-get mm k) (hash-map-find mm k (list)))
(define (multi-map-put! mm k v) (let ((vs (multi-map-get mm k)))
  (hash-map-bind! mm k (cons v vs)))))

(define (new-node data l-node r-node) (vector data l-node r-node))
(define (sentinel) (list))

(define (tree-dfs f acc root)
  (let loop ((node root) (acc acc))
    (match node
      (#(data l-child r-child)
       (loop l-child
         (loop r-child
           (f data acc))))))
  (()) acc)))

```

```
(define (go-left? depth branch)
  (zero? (&& (<< 1 depth) branch))))
```

D.2 Serialization

We then define some wrappers on the standard functions to be able to easily serialize numbers and bytes.

```
(library (sse serialization)
  (export make-bytevector
    bytevector?
    bytevector-length
    string->utf8
    write-u8!    read-u8
    write-u16!   read-u16
    write-u32!   read-u32
    write-u64!   read-u64
    write-u128!  read-u128
    write-bytes! read-bytes)
  (import (sse utils)
    (rnrs base)
    (rnrs control)
    (rnrs bytevectors)
    (rnrs arithmetic bitwise)))

(define endianess% (endianness big))

(define (write-u8! bytes pos val)
  (bytevector-u8-set! bytes pos val)
  (+ pos 1))

(define (write-u16! bytes pos val)
  (bytevector-u16-set! bytes pos val endianess%)
  (+ pos 2))

(define (write-u32! bytes pos val)
  (bytevector-u32-set! bytes pos val endianess%)
  (+ pos 4))

(define (write-u64! bytes pos val)
  (bytevector-u64-set! bytes pos val endianess%)
  (+ pos 8))

(define (write-u128! bytes pos val)
  (let ((v1 (>> val 64))
        (v2 (&& val (- (<< 1 64) 1))))
    (write-u64! bytes (+ pos 0) v1)
    (write-u64! bytes (+ pos 8) v2)
    (+ pos 16)))
```

```

(define (read-u8 bytes pos)
  (values (bytevector-u8-ref bytes pos)
          (+ pos 1)))

(define (read-u16 bytes pos)
  (values (bytevector-u16-ref bytes pos endianess%)
          (+ pos 2)))

(define (read-u32 bytes pos)
  (values (bytevector-u32-ref bytes pos endianess%)
          (+ pos 4)))

(define (read-u64 bytes pos)
  (values (bytevector-u64-ref bytes pos endianess%)
          (+ pos 8)))

(define (read-u128 bytes pos)
  (let ((v1 (read-u64 bytes (+ pos 0)))
        (v2 (read-u64 bytes (+ pos 8))))
    (values (+ (lshift v1 64) v2)
            (+ pos 16)))

(define write-bytes!
  (case-lambda
    (((dest pos src) (write-bytes! dest pos src 0))
     ((dest pos src pos*)
      (let ((len (- (bytevector-length src) pos*)))
        (bytevector-copy! src pos* dest pos len)
        (+ pos len)))))

(define read-bytes
  (case-lambda
    (((src start) (read-bytes src start (bytevector-length src)))
     ((src start stop)
      (let* ((len (- stop start))
             (dest (make-bytevector len)))
        (bytevector-copy! src start dest 0 len)
        (values dest stop))))))

```

D.3 Cryptographic primitives

For the cryptographic operations, we use bindings to GnuTLS.

```

(library (sse crypto)
  (export rng sha3
          aes-gcm-256:init
          aes-gcm-256:keygen
          aes-gcm-256:encrypt
          aes-gcm-256:decrypt)
  (import (rnrs base) (gnutls) (sse serialization)))

```

```

(define (rng n) (gnutls-random random-level/key n))

(define (sha3 bytes) (hash-direct digest/sha3-256 bytes))

(define nonce-length 12)

(define associated-data (make-bytevector 0))

(define (aes-gcm-256:keygen) (rng 32))

(define (aes-gcm-256:init key)
  (make-aead-cipher cipher/aes-256-gcm key))

(define (aes-gcm-256:encrypt cipher ptx)
  (let* ((nonce (rng nonce-length))
         (bytes (aead-cipher-encrypt
                   cipher nonce associated-data 0 ptx)))
    (let ((ctx (make-bytevector (+ nonce-length (bytevector-length bytes)))))
      (write-bytes! ctx 0 nonce)
      (write-bytes! ctx nonce-length bytes)
      ctx)))

(define (aes-gcm-256:decrypt cipher ctx)
  (let ((nonce (read-bytes ctx 0 nonce-length))
        (bytes (read-bytes ctx nonce-length)))
    (aead-cipher-decrypt cipher nonce associated-data 0 bytes)))

```

D.4 The PLOC scheduler

```

(library (sse ploc scheduler)
  (export make-scheduler)
  (import (rnrs base)
          (only (rnrs r5rs) modulo)
          (sse utils))

(define (make-scheduler n)
  (assert (integer? (lg n)))
  (let ((counter 0) (prefixes (iota n)))
    (λ ()
      (set! counter (+ counter 1))
      (map (λ (pfx) (+ pfx (* counter n)))
            prefixes))))

```

D.5 The PLOC scheme

We give here the complete code that allows to run a PLOC client.

```

(library (sse ploc)
  (export make-ploc)
  (import (guile)

```

```

(rnrs base)
(rnrs lists)
(rnrs control)
(sse tree)
(sse utils)
(sse crypto)
(sse serialization)
(sse ploc scheduler)
(srifi :11))

(define (datum value target)
  (vector value target))

(define (datum->value datum) (vector-ref datum 0))
(define (datum->target datum) (vector-ref datum 1))

(define (make-ploc n B V c H value-size read-value write-value!
  ;; Tree dependencies.
  G AE-init AE-keygen AE-encrypt AE-decrypt
  memory-setup memory-read memory-write)

  (assert (u8? c))
  (assert (u64? n))
  (assert (u64? B))
  (assert (u64? V))

  (define datum-size (+ 16 value-size))

  (define (write-datum! bytes pos datum)
    (let* ((pos (write-u128! bytes pos (datum->target datum)))
           (pos (write-value! bytes pos (datum->value datum))))
      pos))

  (define (read-datum bytes pos)
    (let*-values (((target pos) (read-u128 bytes pos))
                  ((value pos) (read-value bytes pos)))
      (values (datum target value)
              pos)))

  (define (data->bytes data)
    (assert (<= (length data) c))
    (let ((bytes (make-u8vector (+ 1 (* c datum-size)))))

      (foldl (λ (pos datum) (write-datum! bytes pos datum))
             (write-u8! bytes 0 (length data))
             data)
      bytes))

  (define (bytes->data bytes)
    (assert (= (+ 1 (* c datum-size)) (bytevector-length bytes))))
```

```

(let-values (((len pos) (read-u8 bytes 0)))
  (assert (<= len c))
  (let loop ((i 0) (pos pos) (data (list)))
    (if (= i len)
        data
        (let-values (((datum pos) (read-datum bytes pos)))
          (loop (+ i 1) pos (cons datum data))))))

(define-values (tree-setup tree-select tree-mutate)
  (make-tree B G data->bytes bytes->data
             AE-init AE-keygen AE-encrypt AE-decrypt
             memory-setup memory-read memory-write))

(define scheduler (make-scheduler n))

(define volumes (hash-map))

(define (get-volume label)
  (hash-map-find volumes label 0))

(define (increment-volume label)
  (let ((volume (get-volume label)))
    (hash-map-bind! volumes label (+ volume 1)))))

(define (get-target key label volume)
  (H key label volume))

(define (make-datum key label value)
  (let* ((volume (get-volume label))
         (target (get-target key label volume)))
    (datum value target)))

(define (get-all-targets key label)
  (map (λ (v) (get-target key label v))
       (iota V)))

(define (tree-data tree)
  (define (index-datum mm datum)
    (match datum
      (#(value target)
        (multi-map-put mm target value))))
  (define (index-data mm data)
    (foldl index-datum mm data))
  (tree-dfs index-data (multi-map) tree))

(define (compact node depth data)
  (define (left-datum? datum)
    (go-left? depth (datum->target datum)))
  (define (direct-data data)
    (split-with left-datum? data))
  (match node

```

```

(#(node-data l-child r-child)
(let*-values
  (((l-data r-data) (direct-data (append node-data data)))
   ((l-child l-data) (compact l-child (+ depth 1) l-data))
   ((r-child r-data) (compact r-child (+ depth 1) r-data)))
   ((node-data rest) (take c (append l-data r-data))))
  (values (new-node node-data l-child r-child) rest)))
  (values (list) data)))

(define (ploc-setup) (tree-setup))

(define (ploc-insert key label value)
  (let* ((data (list (make-datum key label value)))
         (tx (λ (tree)
                (call-with-values (λ () (compact tree 0 data))
                  (λ (tree overflow)
                    (if (null? overflow)
                        tree
                        (error ploc-insert "compaction overflow"
                               label value tree overflow)))))))
    (tree-mutate key (scheduler) tx)
    (increment-volume label)))

(define (ploc-search key label)
  (let* ((targets (get-all-targets key label))
         (subtree (tree-select key targets))
         (values (tree-data subtree)))
    (reverse
      (foldl (λ (acc target)
                 (match (multi-map-get values target)
                   (() acc)
                   ((value) (cons value acc))
                   (values (error ploc-search
                                  "more than one values found"
                                  target values))))
        (list)
        targets))))
  (values ploc-setup ploc-search ploc-insert)))

```

D.6 Secure tree

```

(library (sse tree)
  (export make-tree)
  (import (rnrs base) (sse utils) (sse serialization)))

(define (make-tree D G data->bytes bytes->data
                  AE-init AE-keygen AE-encrypt AE-decrypt
                  memory-setup memory-read memory-write)

(define max-depth D)

```

```

(define (make-address depth target)
  (G depth (&& target (- (<< 1 depth) 1)))))

(define (subtree-addresses branches)
  (let loop ((depth 0) (branches branches))
    (if (or (null? branches) (< max-depth depth))
        (list)
        (let-values
          (((lhs rhs) (split-with (λ (b) (go-left? depth b))
                                    branches)))
           (cons (make-address depth (car branches))
                 (append (loop (+ depth 1) lhs)
                         (loop (+ depth 1) rhs)))))))

(define (word->data AE word)
  (if (null? word)
      (list)
      (bytes->data (AE-decrypt AE word)))))

(define (words->tree AE branches words)
  (let loop ((depth 0) (branches branches) (words words))
    (if (or (null? branches) (< max-depth depth))
        (values (sentinel) words)
        (let*-values
          (((lhs         rhs) (split-with
                                (λ (b) (go-left? depth b))
                                branches))
           ((word      words) (values (car words) (cdr words)))
            ((r-child   words) (loop (+ depth 1) rhs words))
            ((l-child   words) (loop (+ depth 1) lhs words)))
           (values (new-node (word->data AE word)
                            l-child r-child)
                  words))))
        (values (new-node (word->data AE word)
                          l-child r-child)
                words)))))

(define (tree->bindings AE root)
  (let loop ((depth 0) (branch 0) (node root))
    (match node
      (#(data l-child r-child)
       (let ((address (make-address depth branch))
             (word     (AE-encrypt AE (data->bytes data))))
         (cons (cons address word)
               (append (loop (+ depth 1) (+ (<< branch 1) 0) l-child)
                       (loop (+ depth 1) (+ (<< branch 1) 1) r-child)))))))
    (()) (list)))))

(define (tree-setup)
  (memory-setup)
  (AE-keygen))

(define (tree-select key branches)

```

```
(let* ((addresses (subtree-addresses branches))
       (words      (memory-read addresses)))
       (words->tree (AE-init key) branches words)))

(define (tree-mutate key branches transformation)
  (let* ((addresses (subtree-addresses branches))
         (words      (memory-read addresses))
         (AE        (AE-init key))
         (old-tree  (words->tree AE branches words))
         (new-tree  (transformation old-tree)))
    (memory-write (tree->bindings AE new-tree)))))

(values tree-setup tree-select tree-mutate)))
```