

All paths lead to the root

We are on e-print!

The screenshot shows a web page from the Cryptology ePrint Archive. At the top left is the logo, which is a circular arrangement of letters: 'c' at the top, 'a' and 'r' on the bottom right, 't' and 'o' on the bottom left, and 'i' and 'l' on the top right. To the right of the logo is the text "Cryptology ePrint Archive". On the far right of the header are navigation links: "Papers", "Submissions", "About", and a search icon. Below the header, the main content area has a white background. On the left, there's a sidebar with blue and grey abstract shapes. The main content starts with "Paper 2025/1672" and the title "All Paths Lead to the Root". Below the title are the authors: "Théophile Brézot, Cosmian" and "Chloé Hébant, Cosmian". A "Abstract" section follows, containing a detailed description of the research. To the right of the abstract are several sections: "Metadata", "Available format(s)" (with a PDF link), "Category" (with "Cryptographic protocols" highlighted in red), "Publication info" (listing "Preprint"), "Keywords" (listing "SSE STE EMM"), "Contact author(s)" (listing email addresses: "theophile.brezot @ cosmian.com" and "chloe.hebant @ cosmian.com"), and "History" (listing "2025-09-18: approved").

Paper 2025/1672

All Paths Lead to the Root

Théophile Brézot, Cosmian
Chloé Hébant, Cosmian

Abstract

In an attempt to fix the defects of the definition of forward security for Symmetric Searchable Encryption (SSE) schemes, Amjad et al. [2] proposed injection security. This new security property is strictly stronger than most security properties known to date, which makes it particularly challenging to design schemes meeting its requirements. In this work, we show how it is possible to use trees to decorrelate the modification of an index from its effects, hence achieving injection security. In addition to being conceptually simple, our scheme features non-interactive, stateless and mutation-free search operations that allow supporting concurrent readers easily. Finally, the proposed reference implementation is efficient: both Insert and Search operations execute in milliseconds even when operating on an index with up to a million entries and volumes up to a thousand.

Metadata

Available format(s)

[PDF](#)

Category

Cryptographic protocols

Publication info

Preprint.

Keywords

SSE STE EMM

Contact author(s)

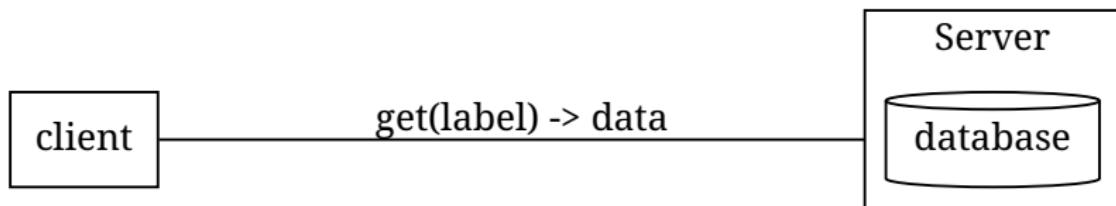
theophile.brezot @ cosmian.com
chloe.hebant @ cosmian.com

History

2025-09-18: approved

Searchable Symmetric Encryption (SSE)

Retrieving data from a server?



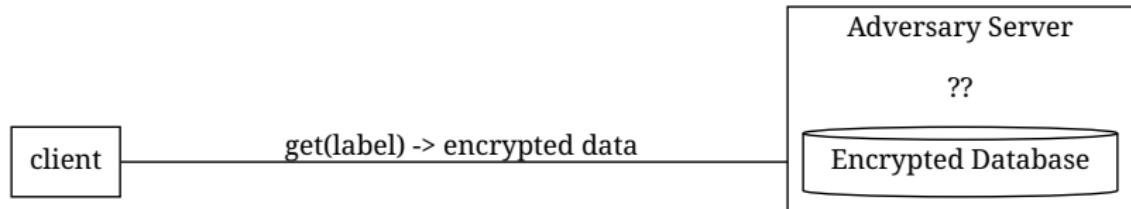
For example: label = "name=DUPONT".

name	surname	birth date	disease
DUPONT	Charles	1971	diabetes
DUPONT	Marie	1988	cancer
...

Retrieving data from a server?

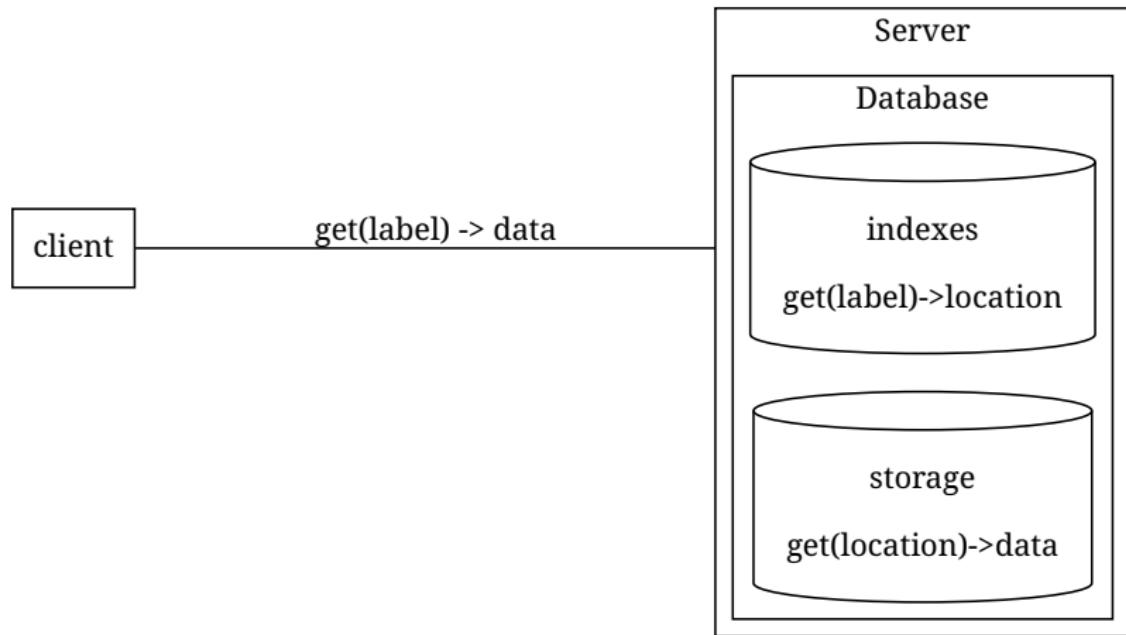
The server knows everything.

Controlling the leakage

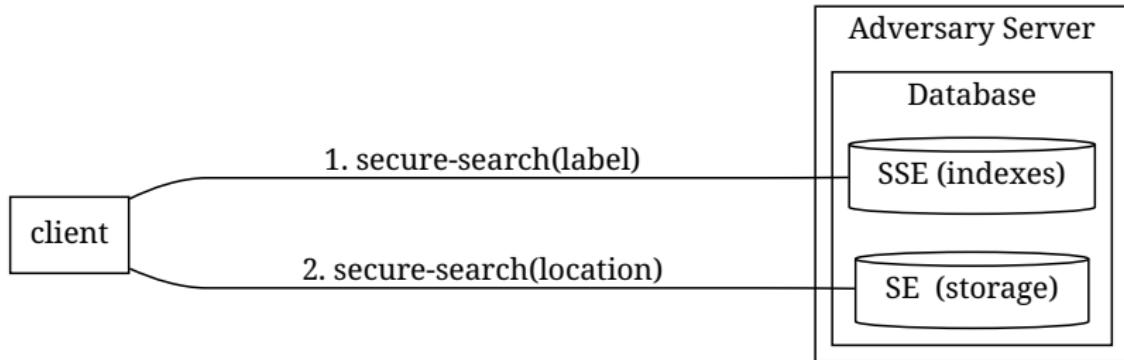


Encrypting the whole database breaks keyword search.

Controlling the leakage



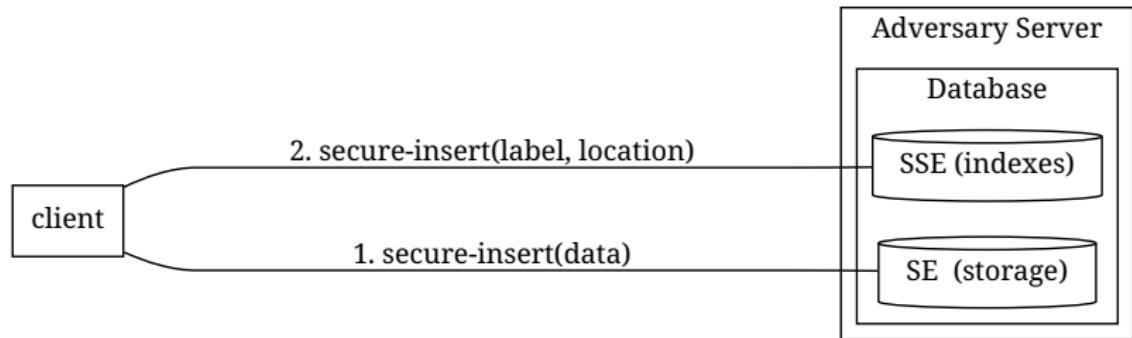
Controlling the leakage



Leaks:

1. $\mathcal{L}_{SSE}(Search, label);$
2. $\mathcal{L}_{SE}(Search, location).$

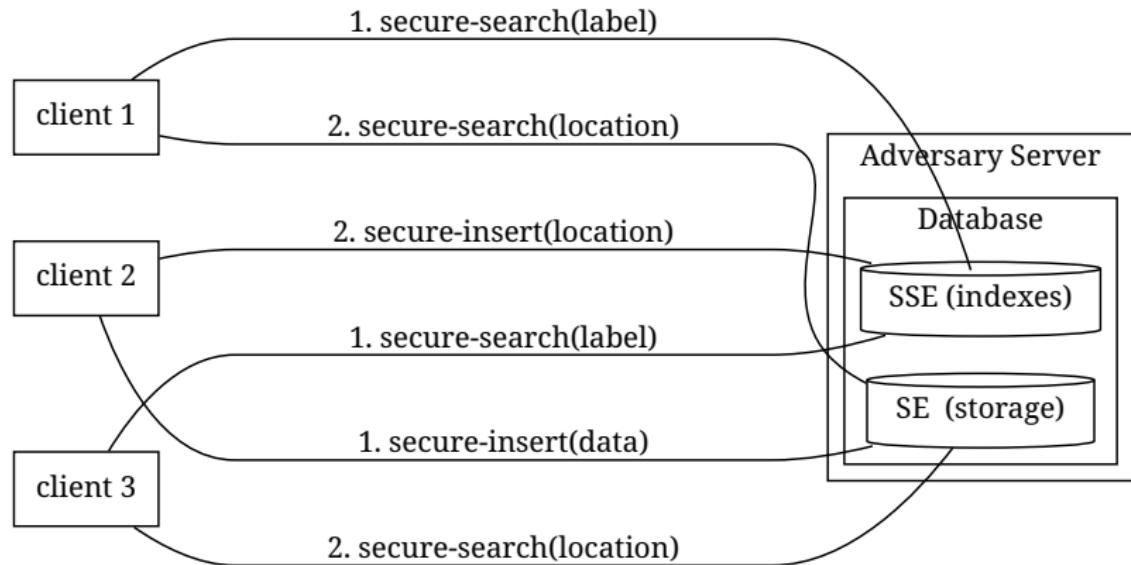
Supporting dynamic states



Leaks:

1. $\mathcal{L}_{SE}(Insert, data)$.
2. $\mathcal{L}_{SSE}(Insert, label, location)$;

Supporting concurrent queries



Structured Encryption

Index

Index ADT:

- ▶ *States:* $\mathbb{L}^{\mathbb{V}^*}$
- ▶ *Operations:*
 - `Search` :: state \rightarrow label \rightarrow Set value
 - `Insert` :: state \rightarrow label \rightarrow value \rightarrow ()

Index STE:

- ▶ *States:* $\mathbb{K}^{\mathbb{L}^{\mathbb{V}^*}}$
- ▶ *Operations:*
 - `Setup` :: () \rightarrow IO key
 - `Search` :: key \rightarrow label \rightarrow IO (Set value)
 - `Insert` :: key \rightarrow label \rightarrow value \rightarrow IO ()

Multi-Map (MM)

MM ADT:

- ▶ *States*: $\mathbb{L}^{\mathbb{V}^*}$
- ▶ *Operations*:
 - `Search` :: state \rightarrow label \rightarrow List value
 - `Insert` :: state \rightarrow label \rightarrow value \rightarrow ()

MM STE:

- ▶ *States*: $\mathbb{K}^{\mathbb{L}^{\mathbb{V}^*}}$
- ▶ *Operations*:
 - `Setup` :: () \rightarrow IO key
 - `Search` :: key \rightarrow label \rightarrow IO (List value)
 - `Insert` :: key \rightarrow label \rightarrow value \rightarrow IO ()

Multi-Map STE to Index STE transformation

```
Index::setup  = MM::setup
Index::insert = MM::insert
Index::search = unique . MM::search
               where
                  unique = fold Set::insert EmptySet
```

Why bother implementing a more constrained
ADT?

Semi-dynamic MM to fully-dynamic MM transformation

Associative ADT:

- ▶ *States:* $\mathbb{L}^{\mathbb{T}(V)}$
- ▶ *Operations:*

Search :: s → l → T v

Mutate :: s → l → (T v → T v) → ()

Semi-dynamic MM to fully-dynamic MM transformation

The fully-dynamic Multi-Map is an associative ADT:

MM::Search s l v ~ Search s l v

MM::Insert s l v ~ Mutate s l (Cons v)

MM::Delete s l v ~ Mutate s l (remove v)

where

remove v [] = []

remove v [v l] = remove v l

remove v [w l] = Cons w (remove v l)

Semi-dynamic MM to fully-dynamic MM transformation

Mutations are composable¹: all we need is to log them!

Journaling Multi-Map ADT:

- ▶ *States:* $\mathbb{L}^{(\mathbb{T}(\mathbb{V})^{\mathbb{T}(\mathbb{V})})^*}$
- ▶ *Operations:*

`Search state -> l -> List (T v -> T v)`

`Insert state -> l -> (T v -> T v) -> ()`

¹More precisely, they form a monoid and can therefore be reduced.

Semi-dynamic MM to fully-dynamic MM transformation

Implementing any (fully-dynamic) associative ADT on top of a (semi-dynamic) multi-map is therefore simple!

```
search s l      = let transformations = MM::search s l
                  in (reduce transformations) T::empty

mutate s l m v = MM::insert s l (m v)
```

PLOC is actually *simple*!

Challenge

MM STE:

- ▶ *States:* $\mathbb{K}^{\mathbb{L}^{V^*}}$
- ▶ *Operations:*
 - `Setup` :: () -> IO key
 - `Search` :: key -> label -> IO (List value)
 - `Insert` :: key -> label -> value -> IO ()

Objectives:

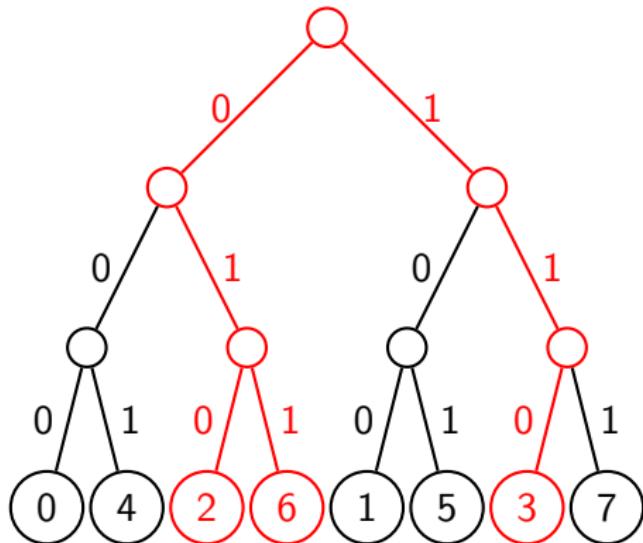
1. Do not leak anything during insertion:
 $\mathcal{L}(Insert, label, value) = \perp$
2. Only leak a (meaningless) UID of the label:
 $\mathcal{L}(Search, label) = sp$

Search

Simply derive the set of target branches:

- ▶ $\text{PRF}(\text{key}, \text{cat}, 0) = 3 \bmod 8 = \text{b...011}$
- ▶ $\text{PRF}(\text{key}, \text{cat}, 1) = 2 \bmod 8 = \text{b...010}$
- ▶ $\text{PRF}(\text{key}, \text{cat}, 2) = 6 \bmod 8 = \text{b...110}$

Mind the endianness!

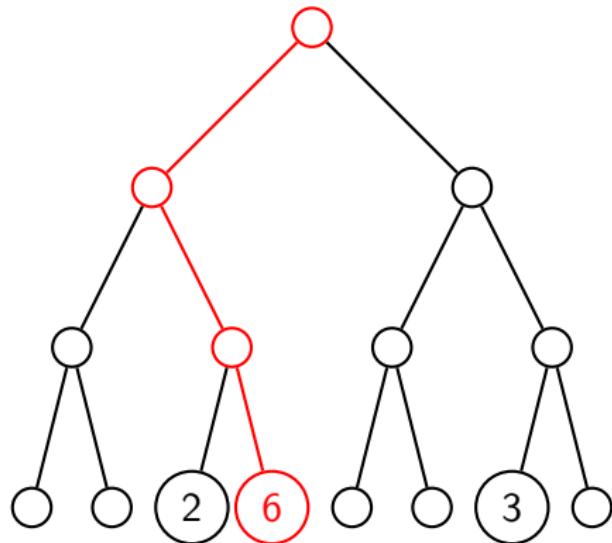


Insert – Datum

```
type Datum = target * value
```

A datum must always be stored on its target branch.

For example with (6 food):



Insertion – Uniform Scheduling

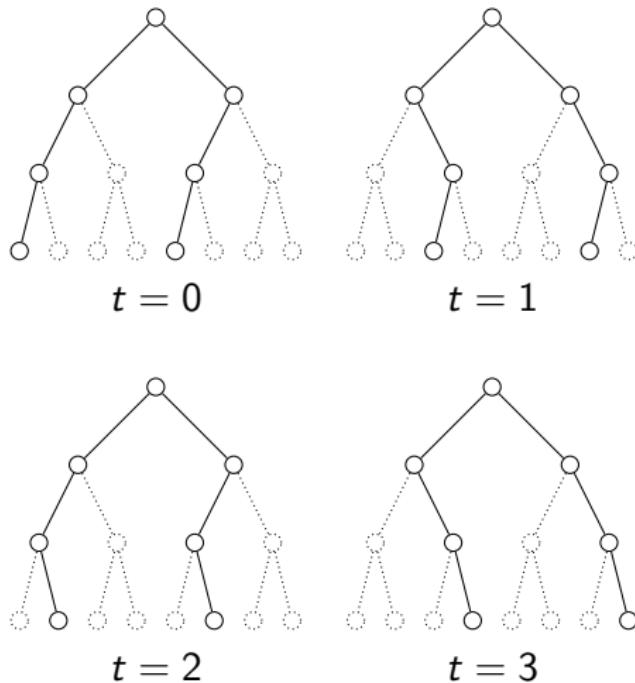
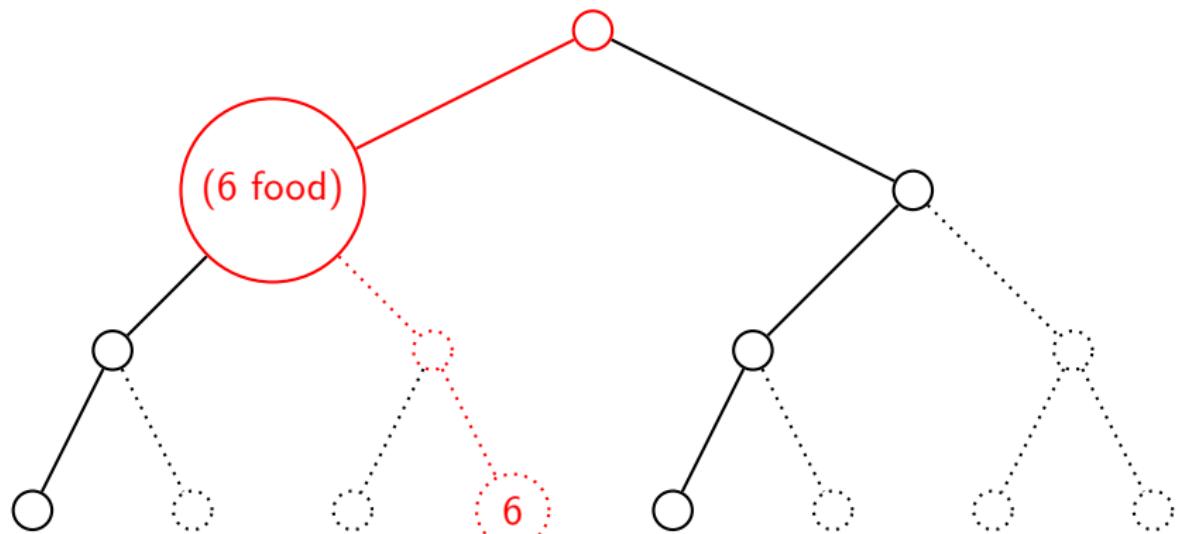
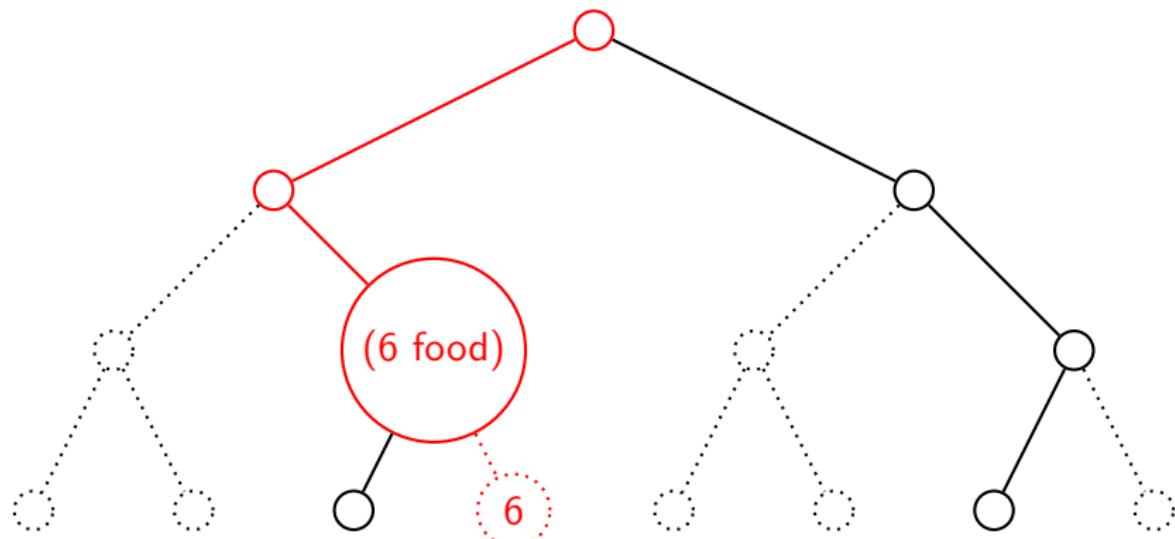


Figure 1: Scheduled subtrees for $N = 8$ and $n = 2$.

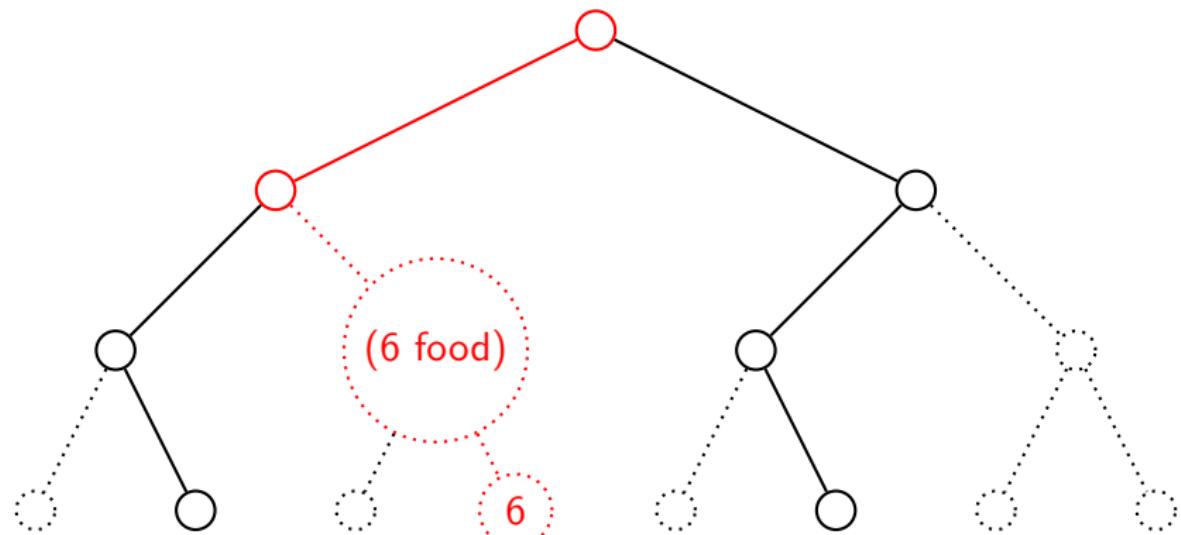
Insert – Compaction 1



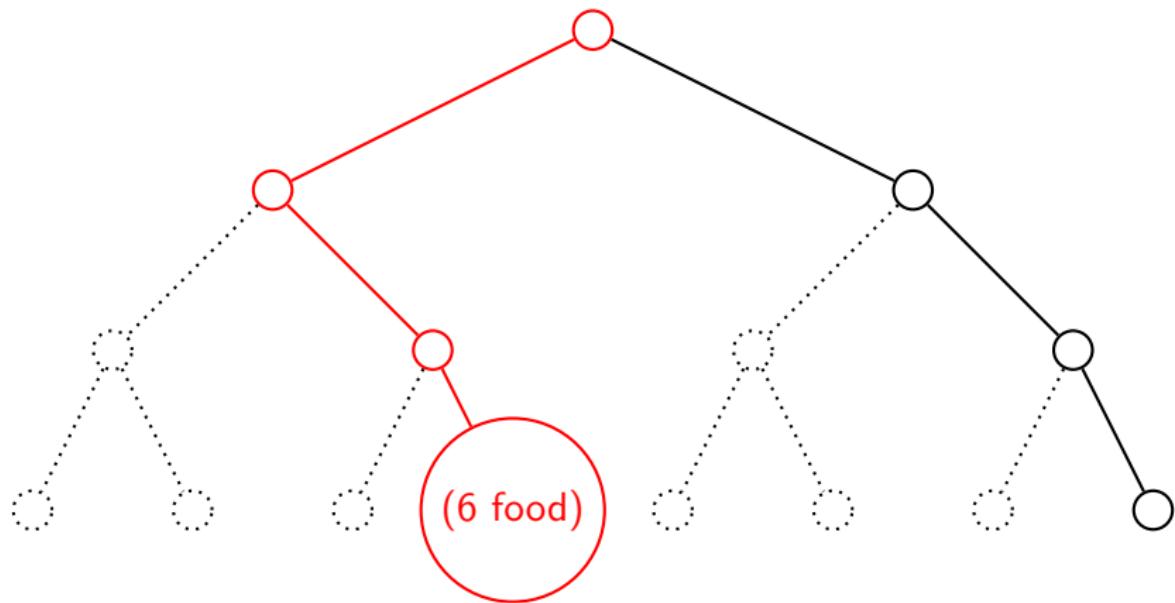
Insert – Compaction 2



Insert – Compaction 3



Insert – Compaction 4



Insert – Compaction

Can the compaction prevent tree overflow?

Insert – Compaction

Can the compaction prevent tree overflow?

yes (but hard to prove)

Conclusion – PLOC is *injection-secure*

- ▶ $\mathcal{L}(\text{Search}, \text{label}) = \{\text{target-branch}\}$
- ▶ $\mathcal{L}(\text{Insert}, \text{label}, \text{value}) = \perp$

Complete decorrelation of mutations and their effects!

Conclusion – PLOC is *efficient*

With a *simple* implementation²:

$n \setminus B$	2^{10}	2^{16}	2^{20}			
16	2.1msec	1.8msec	26msec	3.1msec	0.17sec	3.8msec
64	2.1msec	5.4msec	25msec	11msec	0.12sec	13msec
256	2.0msec	18msec	25msec	34msec	0.12sec	45msec

Table 2: (*Search Insert*) performances in function of n and B for $V = \sqrt{B}$.

²Less than 400LoC!

Conclusion – PLOC is *state-of-the-art*

Only FIX³ achieves the same security properties... but not the same efficiency!

Future works

Can we improve the performance?

- ▶ Search performance is in $O(V)$:
 - ▶ can we store more than one datum per target branch? $\Rightarrow O(\frac{V}{m})$
- ▶ Search bandwidth is in $O(c \lg B)$:
 - ▶ can we reduce the depth by $\lg c$? $\Rightarrow O(\lg B)$

What about concurrency?

Reliance on an synchronized mutable state due to:

- ▶ MM semantics (order)
 - ▶ implement the index directly?
 - ▶ relax progress property?
- ▶ Uniform scheduling (next scheduled branches)
 - ▶ can compaction work with a random scheduling?
 - ▶ relax progress property?

What about concurrency?

Reliance on an synchronized mutable state due to:

- ▶ MM semantics (order) $O(L)$ \Rightarrow bad
 - ▶ ~~implement the index directly?~~
 - ▶ relax progress property + relax target selection
- ▶ Uniform scheduling (next scheduled branches) $O(1)$
 - ▶ ~~can compaction work with a random scheduling?~~
 - ▶ relax progress property?

What about the data-related leakage?

- ▶ Store data directly inside the SSE?
 - ▶ what is the impact on performance?
- ▶ Use an independent scheme with no leakage (ORAM)?
 - ▶ with what performances?
 - ▶ can it be compatible with concurrent queries?

What about the data-related leakage?

- ▶ Store data directly inside the SSE?
 - ▶ what is the impact on performance?
Big due to volume-hiding
- ▶ Use an independent scheme with no leakage (ORAM)?
 - ▶ with what performances?
 - ▶ can it be compatible with concurrent queries?

What about the data-related leakage?

- ▶ Store data directly inside the SSE?
 - ▶ what is the impact on performance?
Big due to volume-hiding
- ▶ Use an independent scheme with no leakage (ORAM)?
 - ▶ with what performances?
OK $O(\log N)$
 - ▶ can it be compatible with concurrent queries?

What about the data-related leakage?

- ▶ Store data directly inside the SSE?
 - ▶ what is the impact on performance?
Big due to volume-hiding
- ▶ Use an independent scheme with no leakage (ORAM)?
 - ▶ with what performances?
OK $O(\log N)$
 - ▶ can it be compatible with concurrent queries?
NO requires a lock

Thanks!